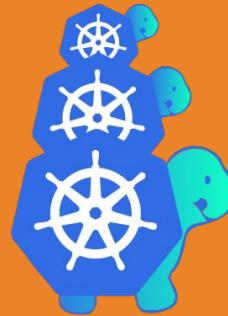


KUBERNETES CLUSTER API



 erkan_erol_

About me

- **Live in Istanbul/Türkiye**
- **Platform Engineer @ Giant Swarm**
- **Netaş -> SAP -> Red Hat**
- **Kubernetes ~ 5 years**
- **Kubernetes operators ~ 3 years**
- **Cluster API ~ 1 year**



Erkan Erol

Disclaimer

- **This presentation contains some icons from [flaticon.com](#)**
- **This talk contains some copy-paste content from Marcel Müller ([twitter.com/MueMarcel](#))'s old talk.**

Scope

For Cluster API users, not for provider maintainers.

The story

How did we get to this point?

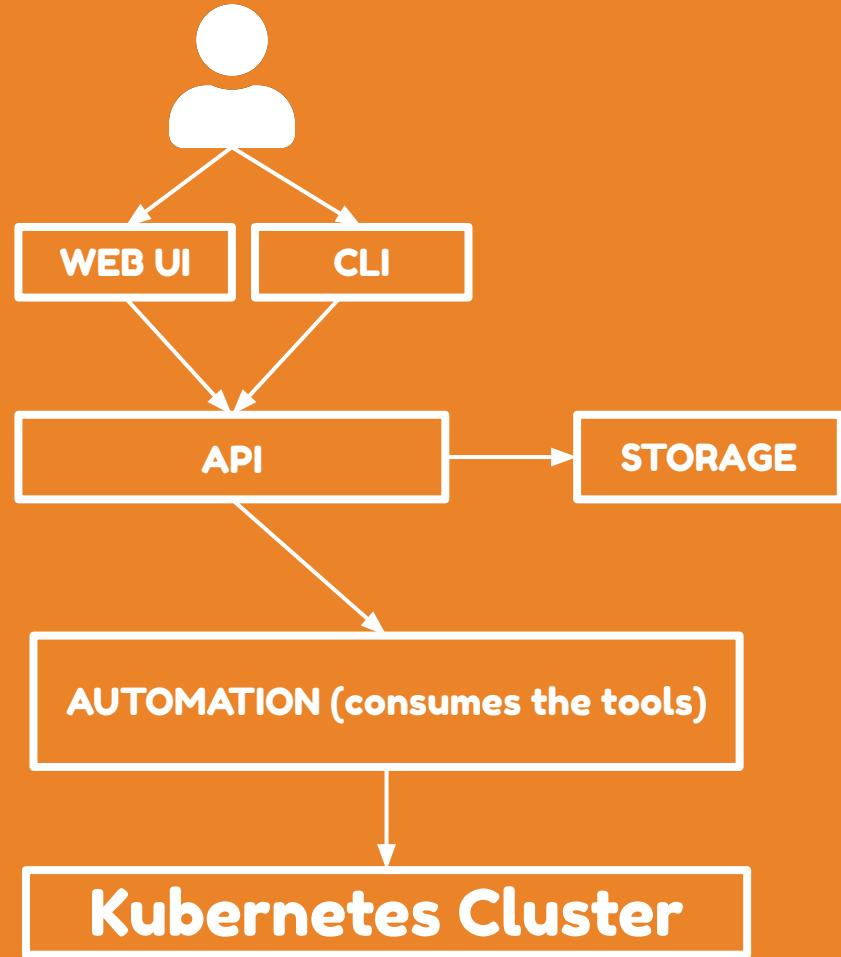
*Everything started
with its birth...*



***Installation was
not easy. There
were many
tools.***



People were implementing systems to manage lifecycle of Kubernetes clusters.



**People started
to use k8s to
deploy
standard
workloads by
using built-in
types.**



DEPLOYMENT



STATEFULSET

Kubernetes

**People liked
Kubernetes
paradigm.**

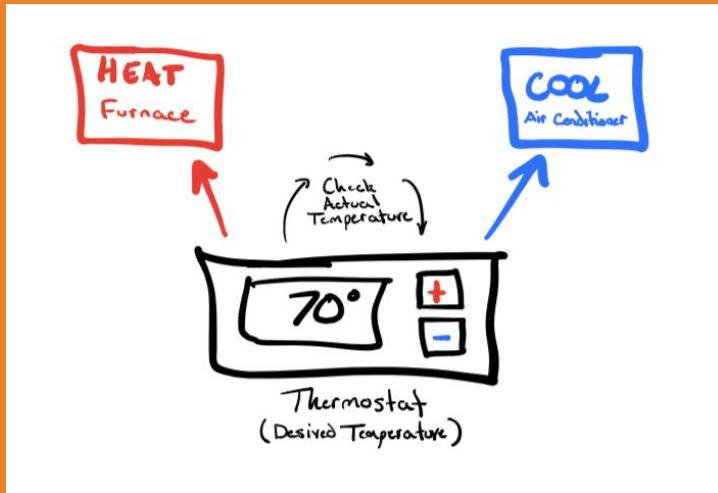


: *I desire this state.*



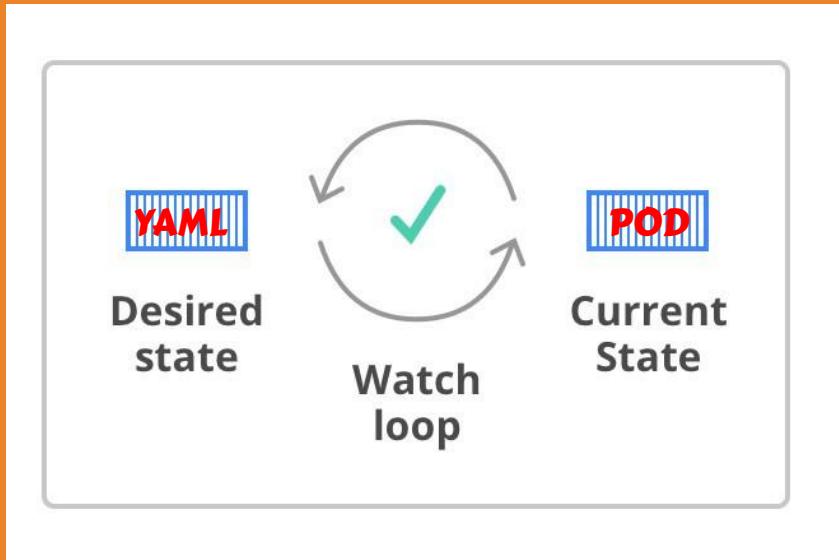
: *Here it is.*

CONTROLLER PATTERN



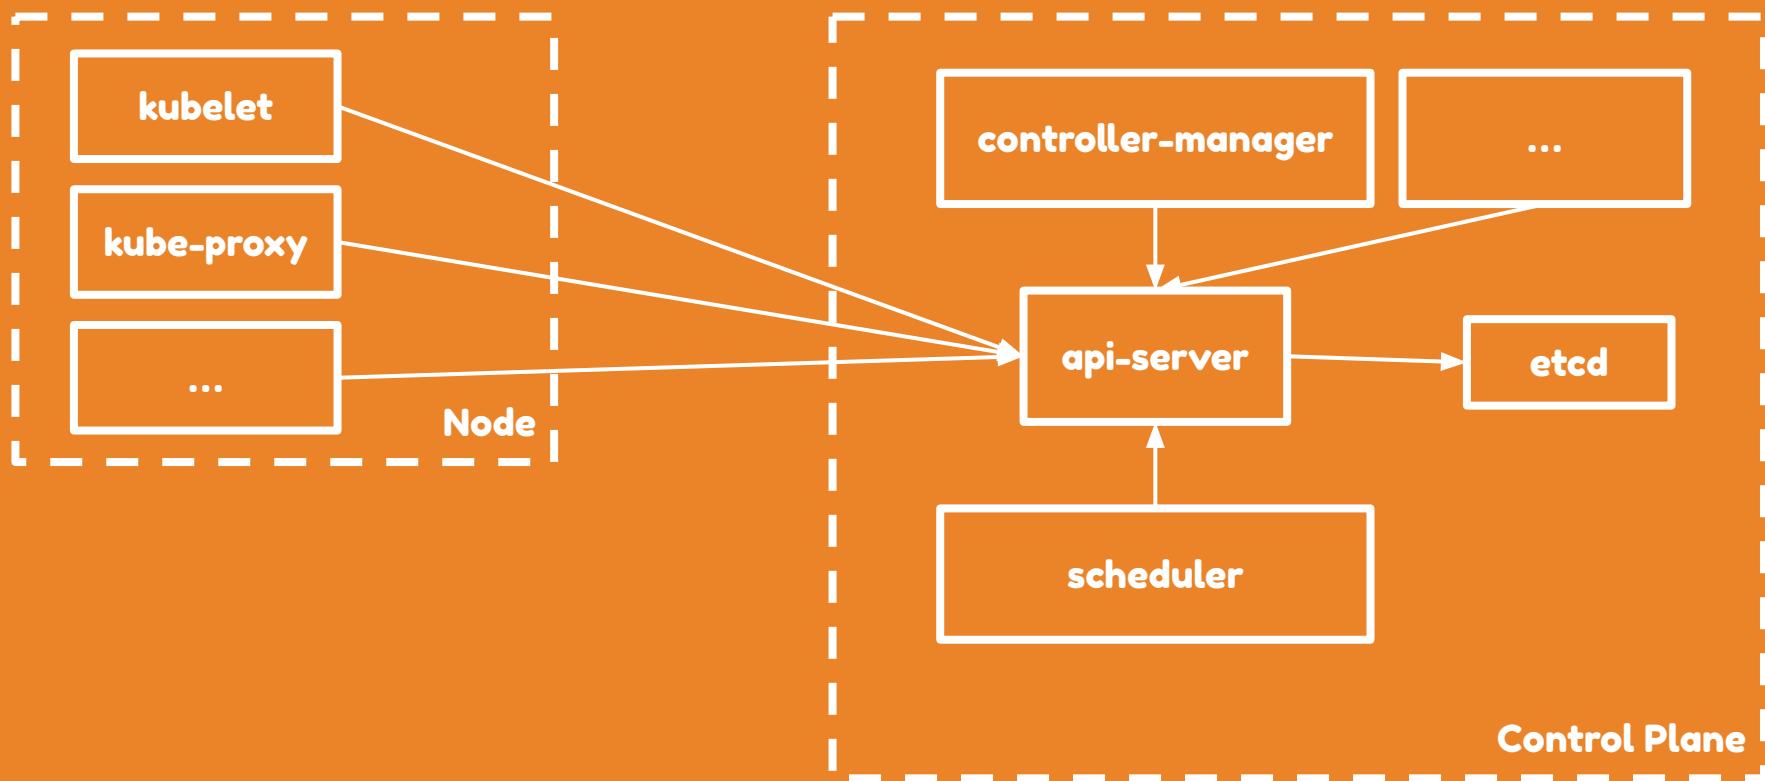
In robotics and automation, a **control loop** is a non-terminating loop that regulates the state of a system.

CONTROLLER PATTERN

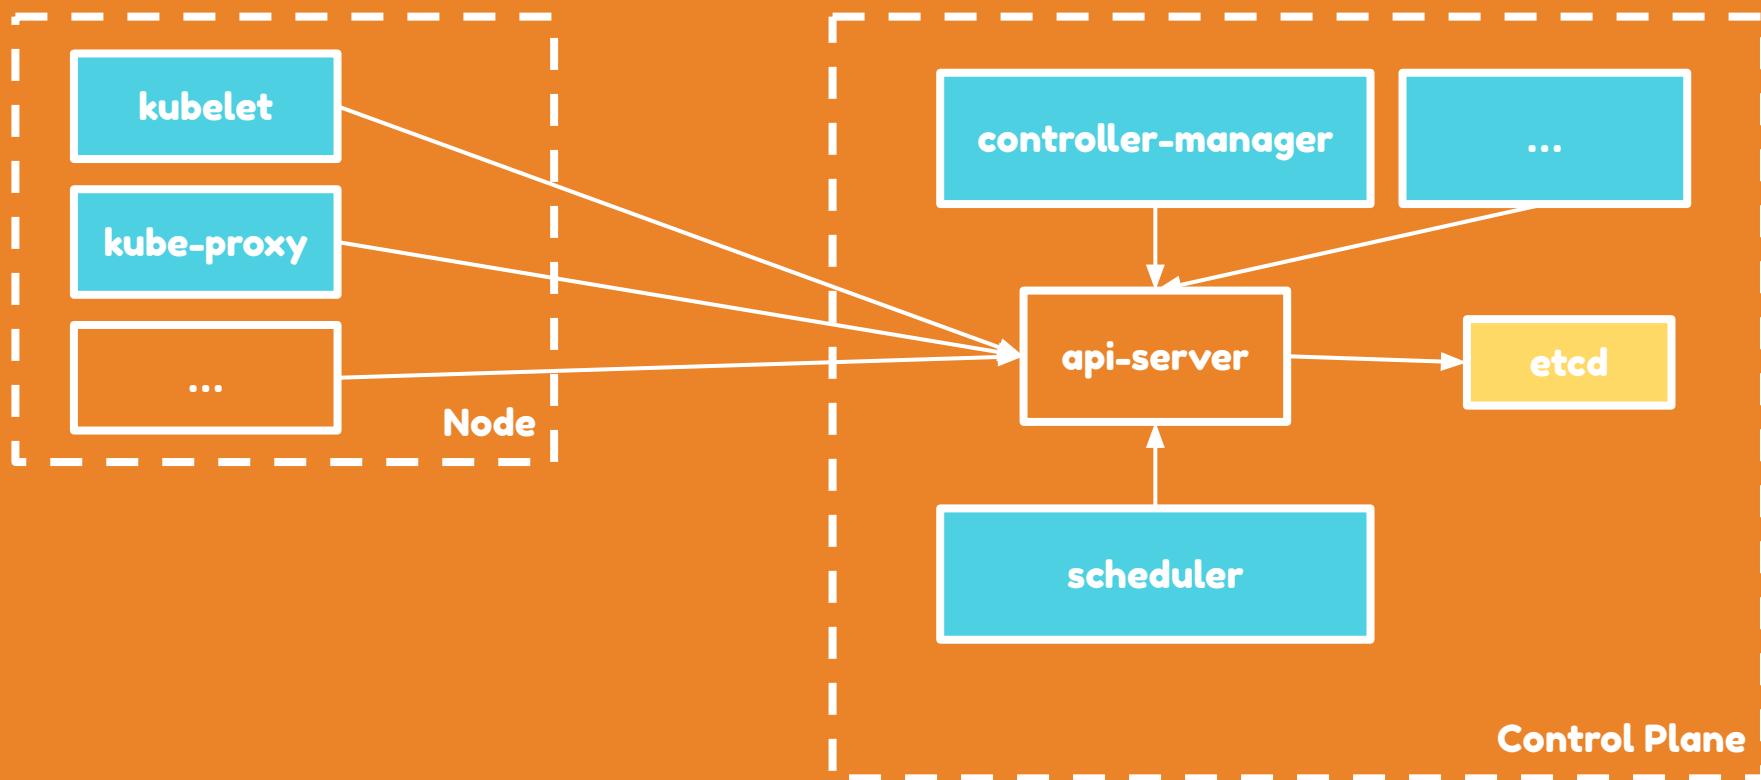


**In Kubernetes,
controllers are control
loops that watch the
state of your cluster, then
make or request changes
where needed.**

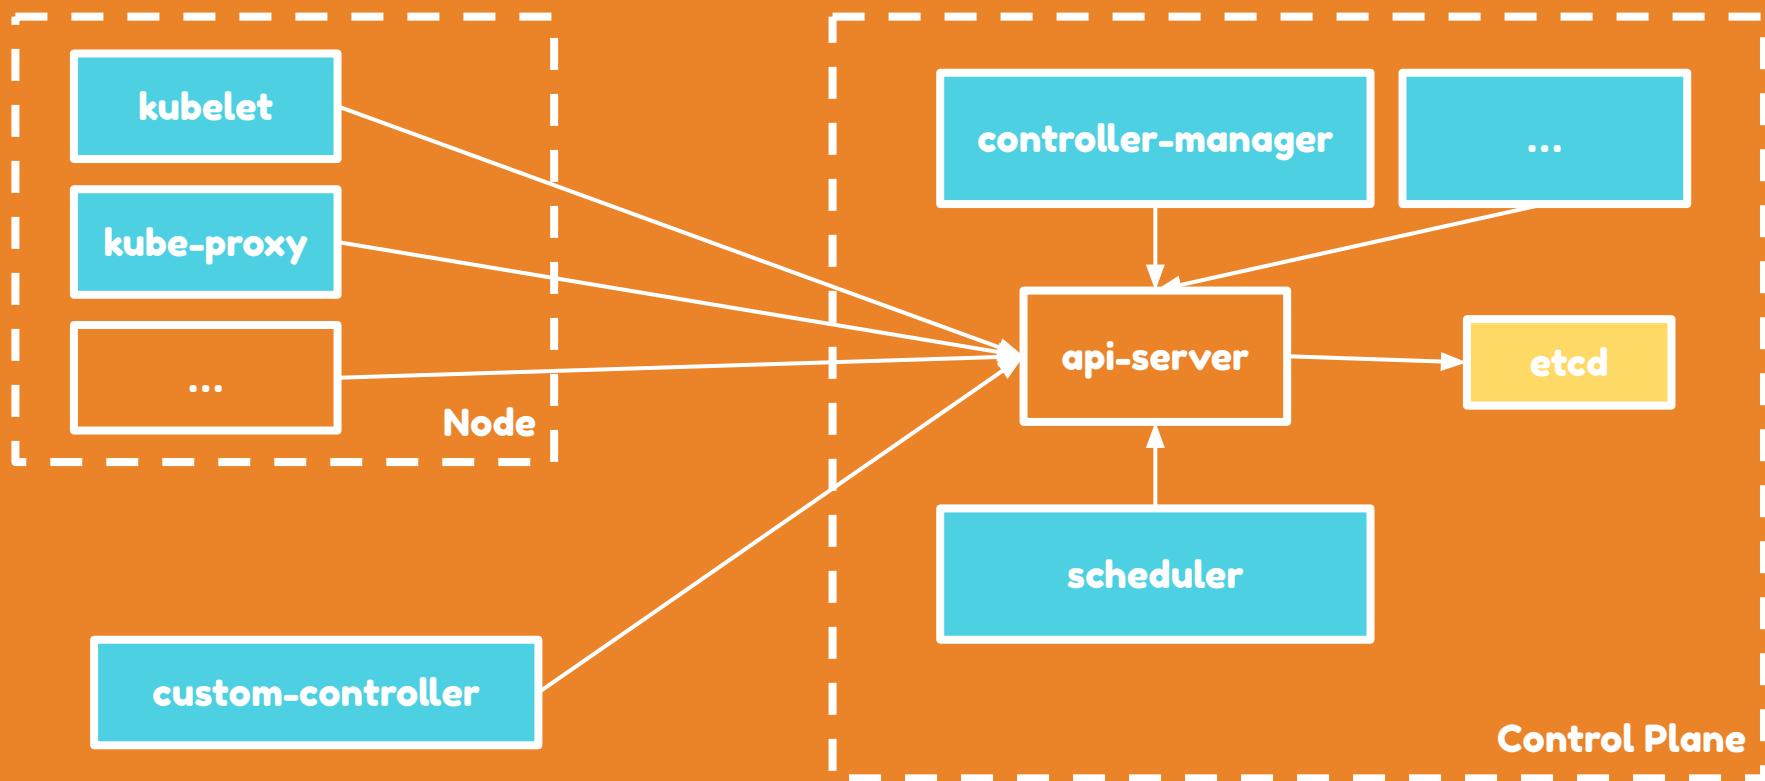
Kubernetes Architecture is actually so simple



storage(etcd) + common api + controllers



Why don't we add yet another controller?



Introducing Operators: Putting Operational Knowledge into Software

November 3, 2016 · By Brandon Philips

A Site Reliability Engineer (SRE) is a person that operates an application by writing software. They are an engineer, a developer, who knows how to develop software specifically for a particular application domain. The resulting piece of software has an application's operational domain knowledge programmed into it.

Operator Pattern

- **aims to capture the key aim of a human operator**
 - **how the system ought to behave**
 - **how to deploy it**
 - **how to react if there are problems**
- **k8s operator = one or more controllers for a specific thing**
- **e.g. MySql Operator, Prometheus Operator etc.**

People
♥ liked ♥
this
more



: I desire this

```
apiVersion: monitoring.coreos.com/v1
kind: Prometheus
metadata:
  name: example
spec:
  replicas: 2
  alerting:
    alertmanagers:
    - namespace: default
      name: alertmanager-example
      port: web
  serviceMonitorSelector:
    matchLabels:
      team: frontend
  ruleSelector:
    matchLabels:
      role: alert-rules
  prometheus: example
```



prometheus-operator

: Here it is.



**Not only k8s
resources
but also
external
resources
too.**



: I desire this



```
apiVersion: ec2.aws.crossplane.io/v1beta1
kind: VPC
metadata:
  name: sample-vpc
spec:
  forProvider:
    region: us-east-1
    cidrBlock: 10.0.0.0/16
    enableDnsSupport: true
    enableDnsHostNames: true
    instanceTenancy: default
    providerConfigRef:
      name: example
```

: Here it is.





WHY DON'T WE USE OPERATORS TO MANAGE KUBERNETES CLUSTERS?



LET'S STANDARDIZE!

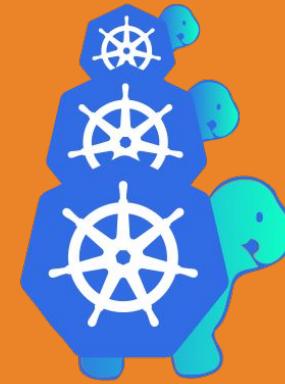
Kubernetes Cluster API

Cluster API is a Kubernetes sub-project focused on providing declarative APIs and tooling to simplify provisioning, upgrading, and operating multiple Kubernetes clusters.

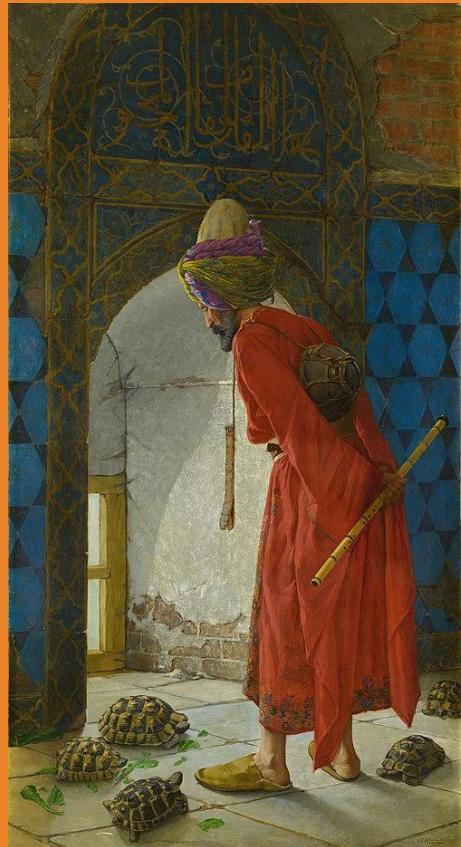
Started by the Kubernetes Special Interest Group (SIG) Cluster Lifecycle, the Cluster API project uses Kubernetes-style APIs and patterns to automate cluster lifecycle management for platform operators. The supporting infrastructure, like virtual machines, networks, load balancers, and VPCs, as well as the Kubernetes cluster configuration are all defined in the same way that application developers operate deploying and managing their workloads. This enables consistent and repeatable cluster deployments across a wide variety of infrastructure environments.



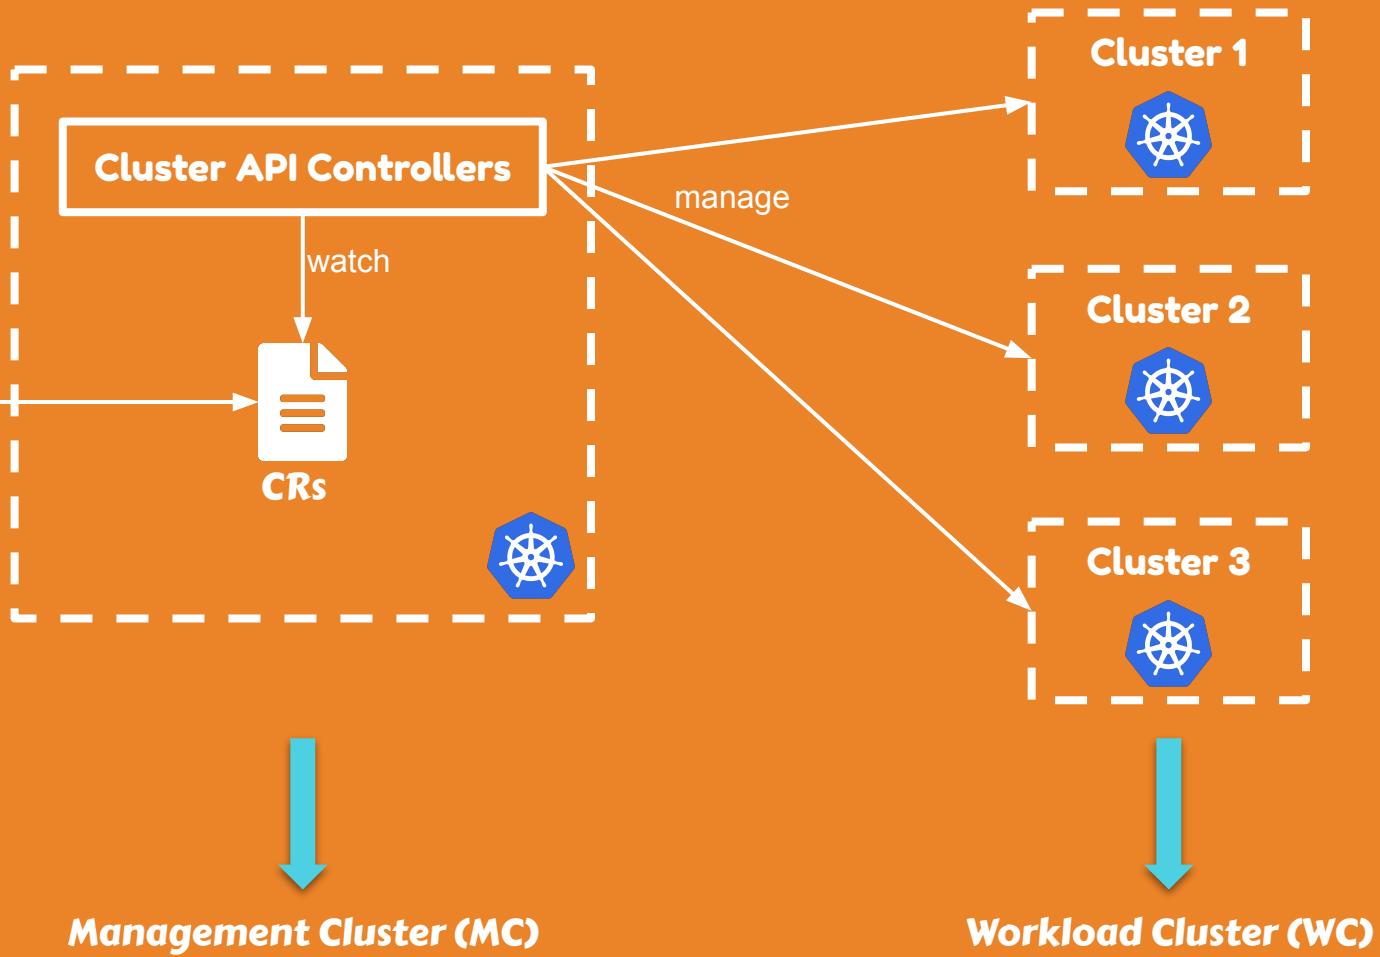
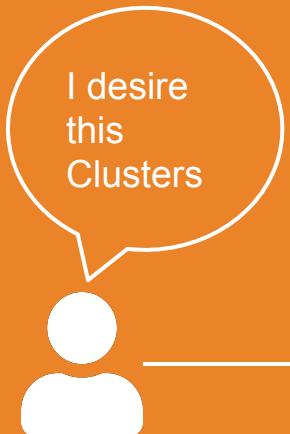
And it was born



Cluster API Fundamentals



The Tortoise Trainer





To understand the solution,

*Let's think about the problem at
first.*



What do we need to create a k8s cluster?

I. We need infra resources



Machines / VMs



Networking



Storage



Infrastructure Provider

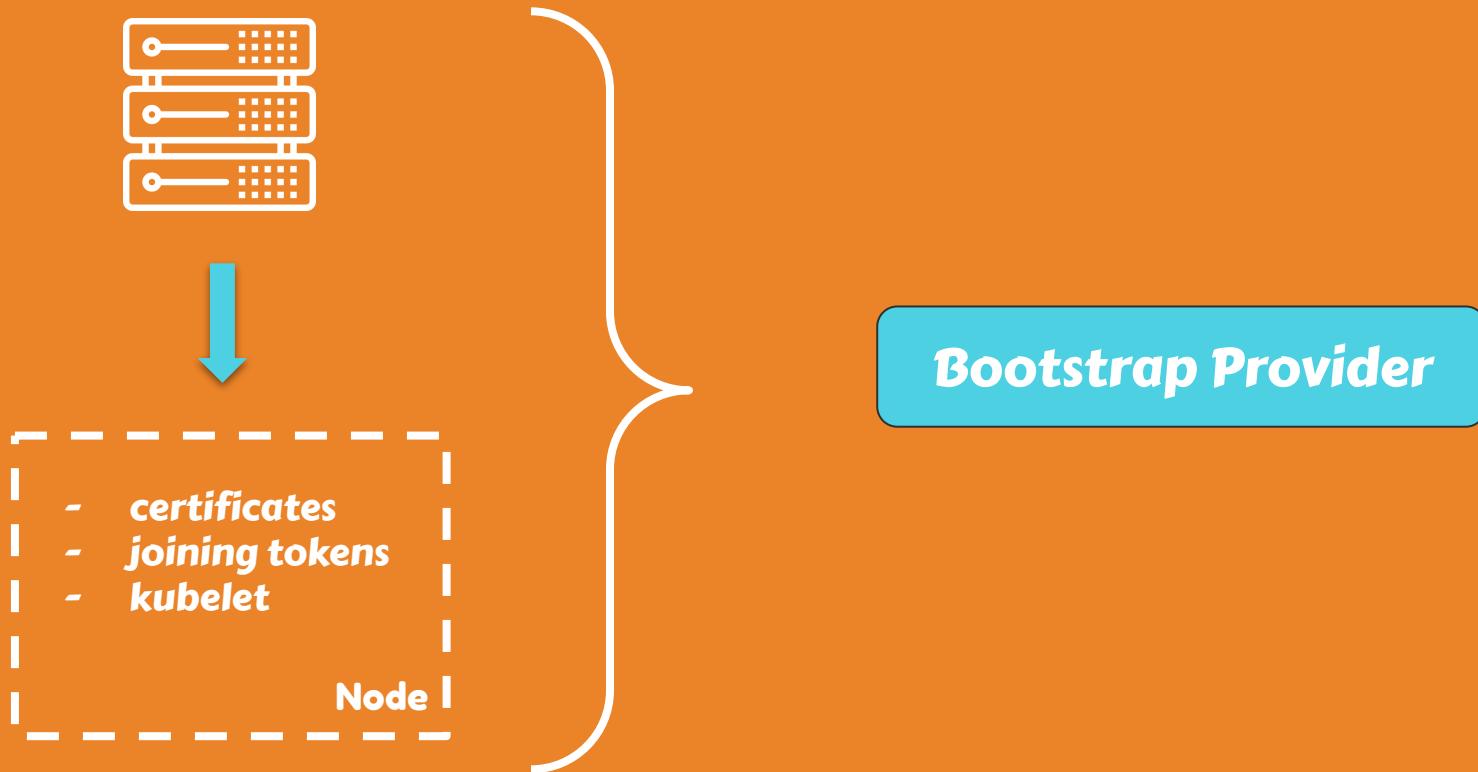
Current Infrastructure Providers

- **AWS**
- **Azure**
- **Azure Stack HCI**
- **BYOH**
- **CloudStack**
- **DigitalOcean**
- **Equinix Metal**
- **GCP**
- **Hetzner**
- **Outscale**
- **IBM Cloud**
- **♥KubeVirt♥**
- **MAAS**
- **Metal3**
- **Microvm**
- **Nested**
- **Nutanix**
- **OpenStack**
- **OCI**
- **Sidero**
- **Tencent Cloud**
- **vcluster**
- **Virtink**
- **VMware Cloud Director**
- **vSphere**



We will use it as an example
in the upcoming slides

2. We need to convert machines to k8s nodes

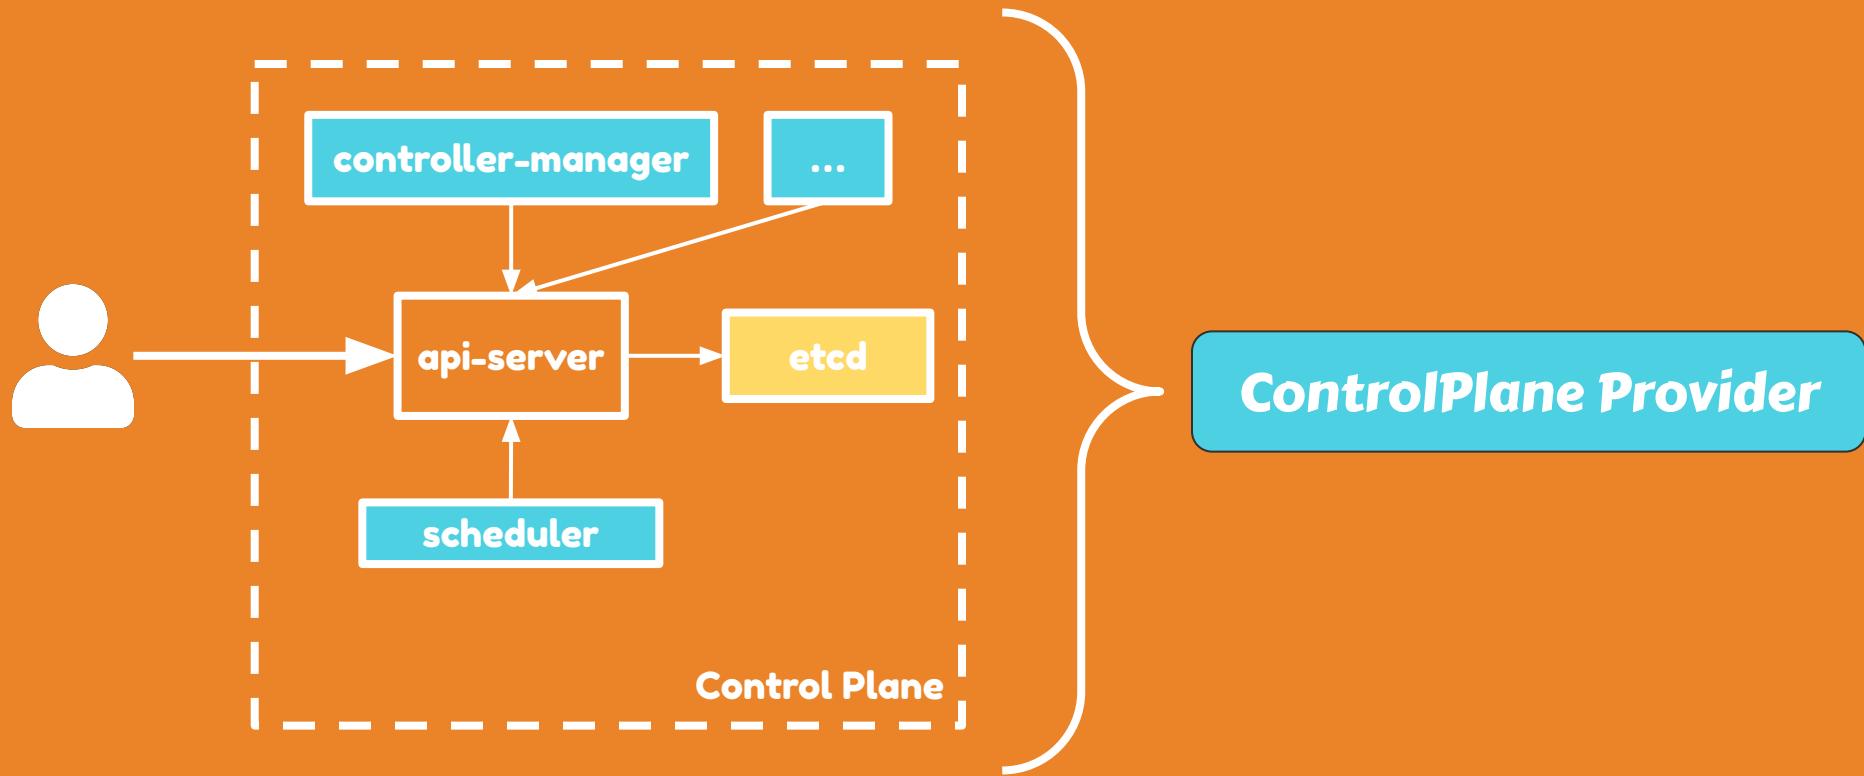


Current Bootstrap Providers

- **EKS**
- **Kubeadm** —————→ **We will talk about only this one**
- **Talos**



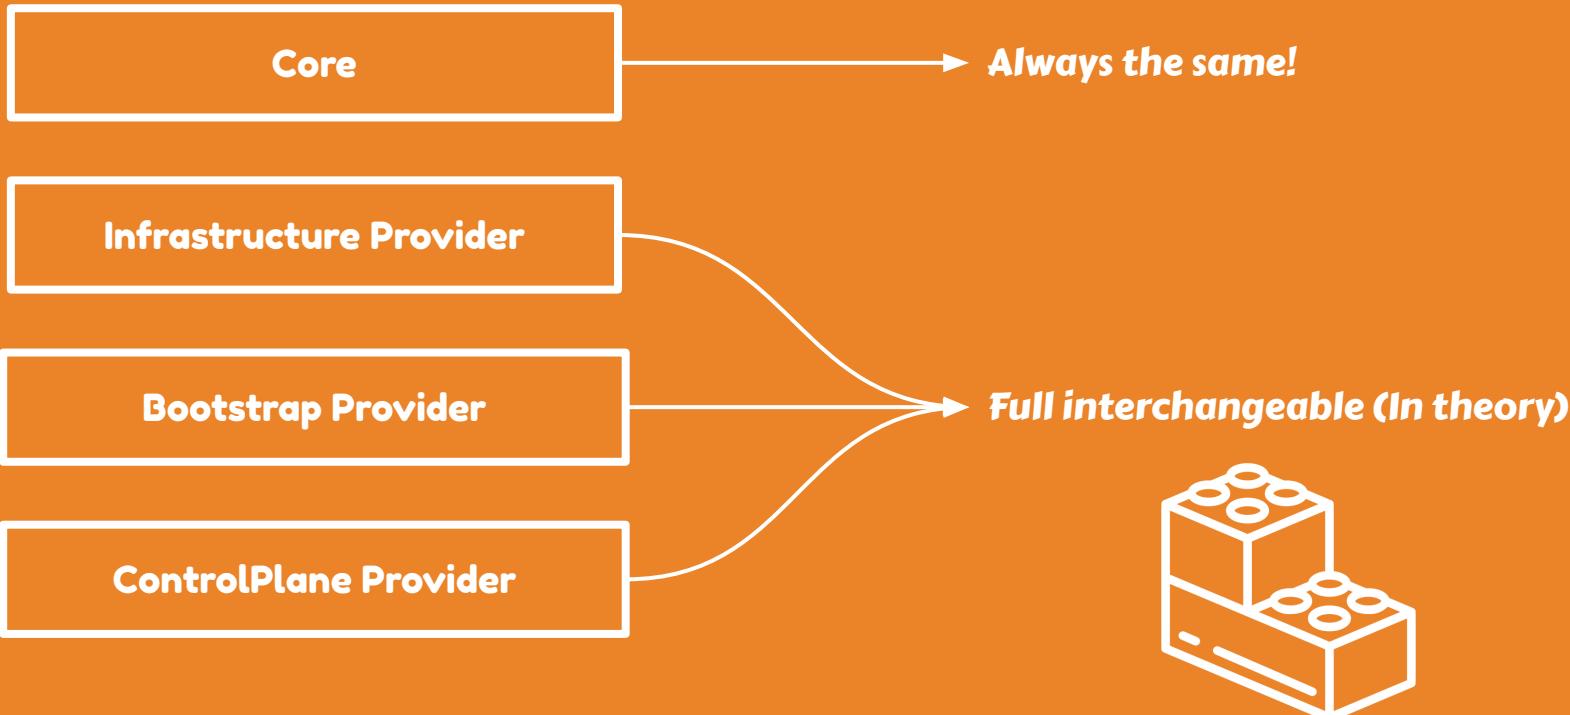
3. We need a control-plane to join our nodes



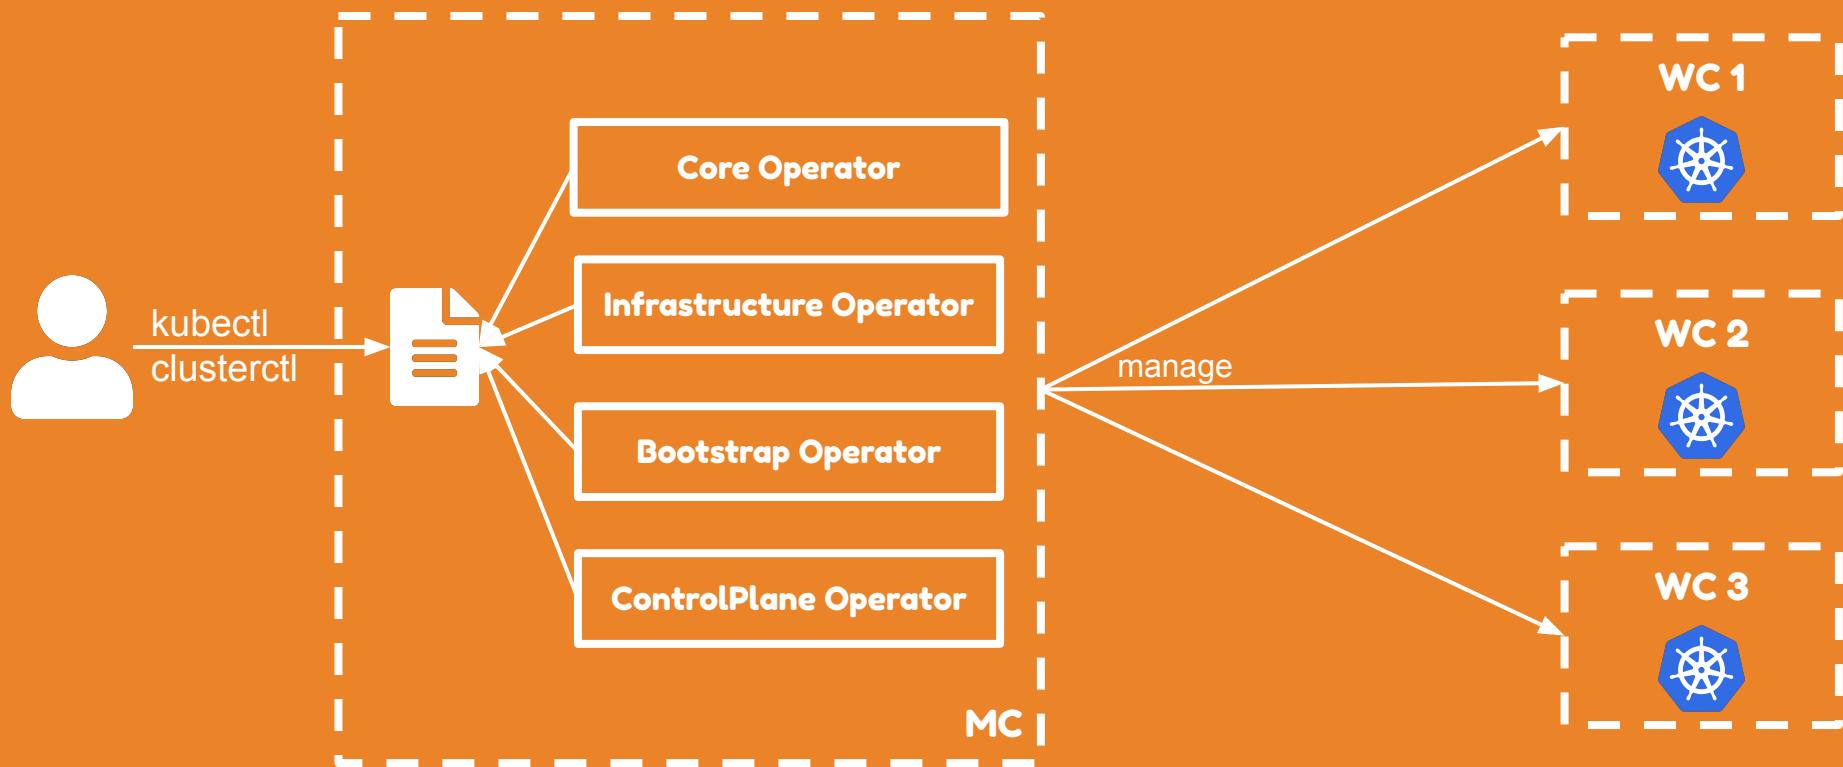
ControlPlane Provider Types

- **Self-provisioned**
 - In WC, controlled by Cluster API
 - e.g. *KubeadmControlPlane*
- We will talk about only this one
- **Pod-based**
 - In external cluster, pod based, exposed by a k8s service
- **External or Managed**
 - controlled by another system like GKE, AKS, EKS, etc.

Cluster API components are interchangeable!



BIG PICTURE



clusterctl

- **The clusterctl CLI tool handles the lifecycle of a Cluster API management cluster.**
- **It automates fetching the YAML files defining provider components and installing them.**

- `clusterctl init` Initialize a management cluster.
- `clusterctl upgrade plan` Provide a list of recommended target versions for upgrading Cluster API providers in a management cluster.
- `clusterctl upgrade apply` Apply new versions of Cluster API core and providers in a management cluster.
- `clusterctl delete` Delete one or more providers from the management cluster.
- `clusterctl generate cluster` Generate templates for creating workload clusters.
- `clusterctl generate yaml` Process yaml using clusterctl's yaml processor.
- `clusterctl get kubeconfig` Gets the kubeconfig file for accessing a workload cluster.
- `clusterctl move` Move Cluster API objects and all their dependencies between management clusters.
- `clusterctl alpha rollout` Manages the rollout of Cluster API resources. For example: MachineDeployments.

clusterctl init example

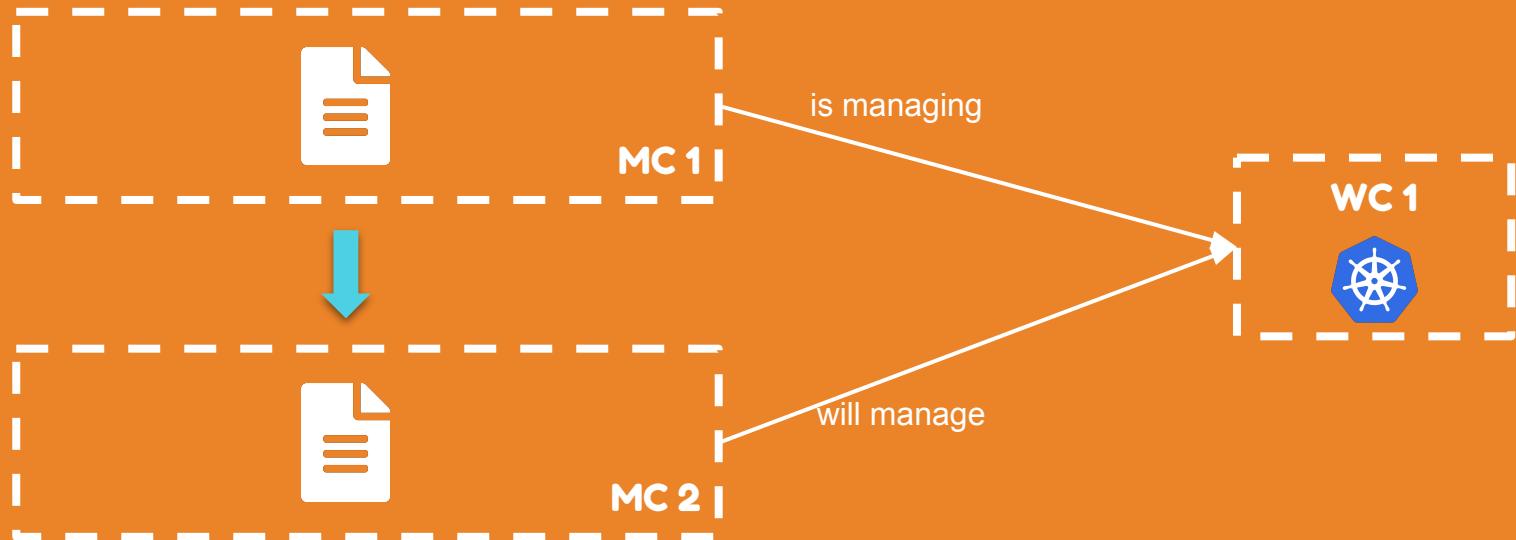
```
→ clusterctl init --infrastructure openstack
Fetching providers
Installing cert-manager Version="v1.7.2"
Waiting for cert-manager to be available...
Installing Provider="cluster-api" Version="v1.2.4" TargetNamespace="capi-system"
Installing Provider="bootstrap-kubeadm" Version="v1.2.4" TargetNamespace="capi-kubeadm-bootstrap-system"
Installing Provider="control-plane-kubeadm" Version="v1.2.4" TargetNamespace="capi-kubeadm-control-plane-system"
Installing Provider="infrastructure-openstack" Version="v0.6.3" TargetNamespace="capo-system"

Your management cluster has been initialized successfully!
```

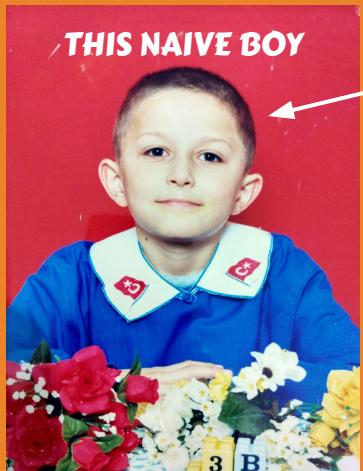
You can now create your first workload cluster by running the following:

```
clusterctl generate cluster [name] --kubernetes-version [version] | kubectl apply -f -
```

clusterctl move



DREAM vs REALITY



November 13th, 2021 ▾

 **Erkan Erol** 7:29 PM

Hi all. I have a high level understanding about cluster-api and I am trying to understand the rationale behind cluster-api architecture/design. I didn't get something.

Before jumping into details, I was under such an impression:

- There are some common types coming from cluster-api core like Cluster, ControlPlane, Machine etc.
- End users are responsible for creating these common types. They have options to pick a specific provider.
- Controllers of providers are watching these common types and they create their own types accordingly and provide the necessary infra.

like the mechanism between PersistentVolumeClaim, StorageClass, PersistentVolume etc.

After checking the examples, my understanding is that users have to create the resources for a specific type like `Metal3Cluster` and put a reference in the `Cluster` obj.

My question: Why? Why don't we have common specs for all providers? Is there a resource about it so that I can check it? (edited)

 4 replies Last reply 11 months ago

DREAM vs REALITY

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
metadata:
  name: cluster-spec-in-my-dream
spec:
  infra: OpenStack
  bootstrap: Kubeadm
  controlPlane: GKE
  myGenericSpec:
    workerNodeReplicas: 3
    controlPlaneReplicas: 3
```

There is no one or few common CRs for all providers.

```
→ helm get manifest demo1-cluster |grep '^kind' -B 1
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
--
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
kind: OpenStackCluster
--
apiVersion: controlplane.cluster.x-k8s.io/v1beta1
kind: KubeadmControlPlane
--
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmConfigTemplate
--
apiVersion: cluster.x-k8s.io/v1beta1
kind: MachineDeployment
--
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
kind: OpenStackMachineTemplate
--
apiVersion: cluster.x-k8s.io/v1beta1
kind: MachineHealthCheck
--
apiVersion: addons.cluster.x-k8s.io/v1beta1
kind: ClusterResourceSet
```

Each provider has its own CRs.



Deep Dive

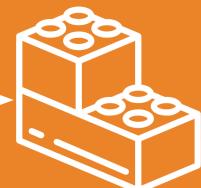


Deep Dive Stack

- **Core Operator**
- **Openstack as infra provider**
- **Kubeadm as bootstrap provider**
- **KubeadmControlPlane as controlplane provider**



**These are just examples to understand the concepts.
Don't forget**



I want to create a cluster with Cluster API.

How can I do it?

kind: Cluster

- **Managed by “capi-controller-manager” (Core Operator)**
- **Main CR**
- **Give reference to infra provider**
- **Give reference to controlplane provider (Optional)**

kind: Cluster : Spec



As a Cluster API user

I want to have a cluster

With this control plane provider

On this infra provider

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster

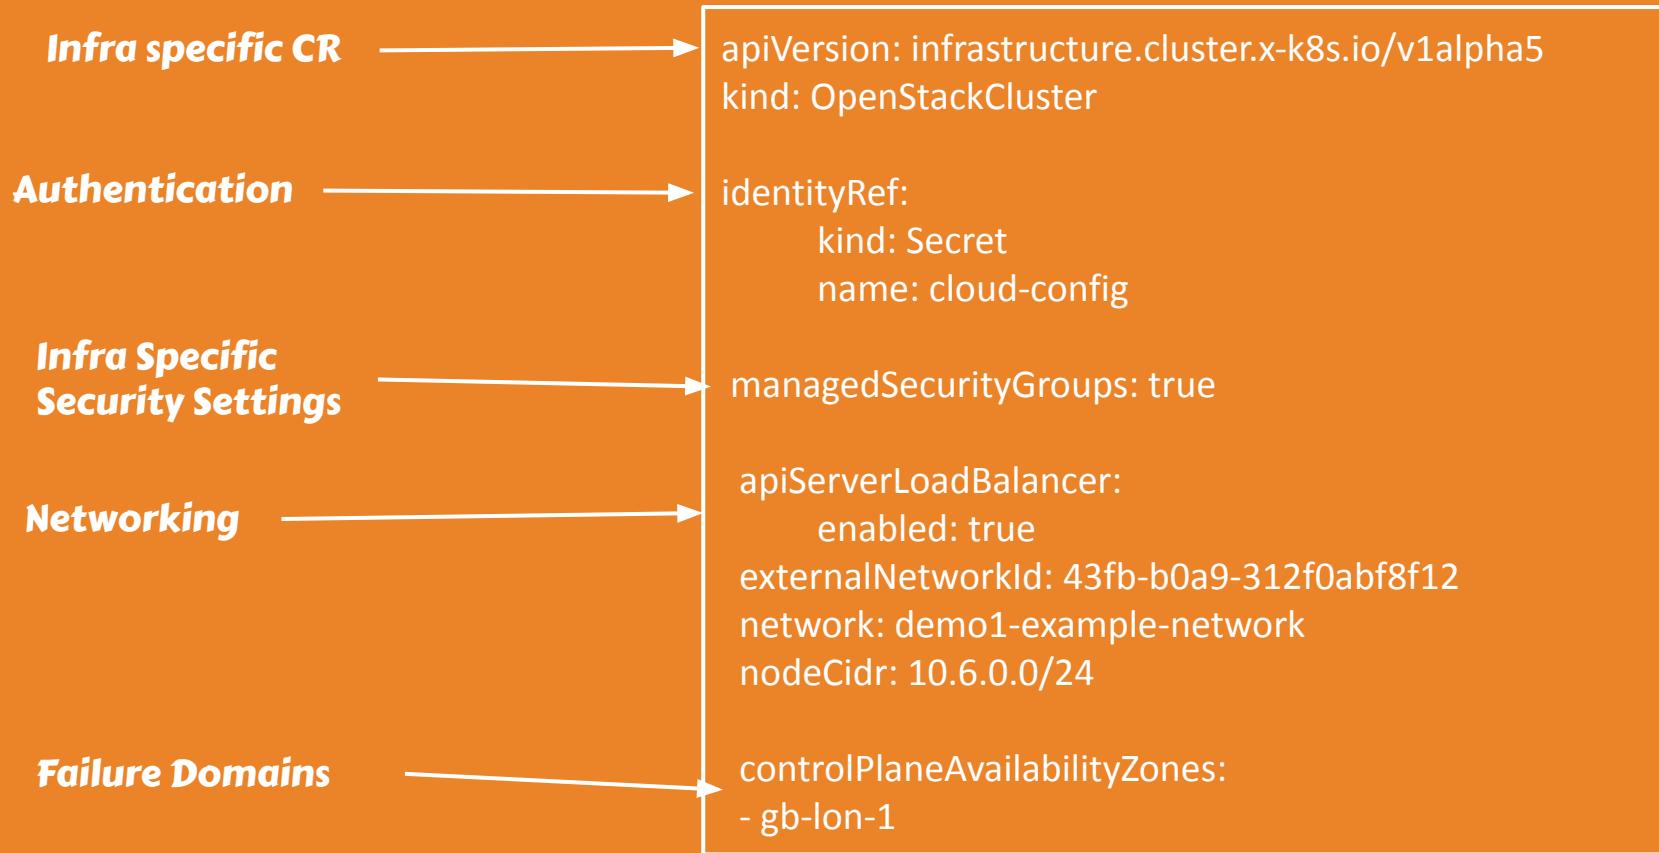
spec:
...
controlPlaneRef:
    apiVersion: controlplane.cluster.x-k8s.io/v1beta1
    kind: KubeADMControlPlane
    name: demo1
    namespace: org-multi-project

infrastructureRef:
    apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
    kind: OpenStackCluster
    name: demo1
    namespace: org-multi-project
```

kind: InfraCluster

- **e.g. OpenStackCluster**
- **Managed by “capo-controller-manager” (Infra Operator)**
- **CR to declare common/primary infra resources for the cluster**
 - **Project / Account in infra provider**
 - **Credentials to access infra provider**
 - **VPC / Network configuration**
 - **Region / Availability Zone / Failure Domains**

kind: InfraCluster : Spec



**With Cluster +
InfraCluster, we created
some basics like project,
vpc, security groups, LB
for k8s api etc. but we
still don't have any
machine!**



kind: Machine

- **Managed by “capi-controller-manager” (Core Operator)**
- **Give reference to infra provider**
- **Give reference to bootstrap provider (Optional)**

kind: Machine : Spec



As a Cluster API user

I want to have a machine

On this infra provider

With this bootstrap configuration

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: Machine

spec:
  infrastructureRef:
    apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
    kind: OpenStackMachine
    name: demo1-default-91da4cd9-m7gn2
    namespace: org-multi-project
    uid: d7ac11dd-d41c-40a4-9a4b-ab2fe5dd71ad

  bootstrap:
    configRef:
      apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
      kind: KubeADMConfig
      name: demo1-region1-8342e0a3-lwlqx
      namespace: org-multi-project
      uid: 5a31ddc4-92e3-4bc9-9af3-d7827e6366fc
      dataSecretName: demo1-region1-8342e0a3-lwlqx
```

kind: InfraMachine

- **e.g. OpenStackMachine**
- **Managed by “capo-controller-manager” (Infra Operator)**
- **CR to declare machine spec**
 - **Image**
 - **Size / Flavor**
 - **Network / Subnet**
 - **Volumes**
 - **...**

kind: InfraMachine : Spec

Infra specific CR



Authentication



Machine Spec

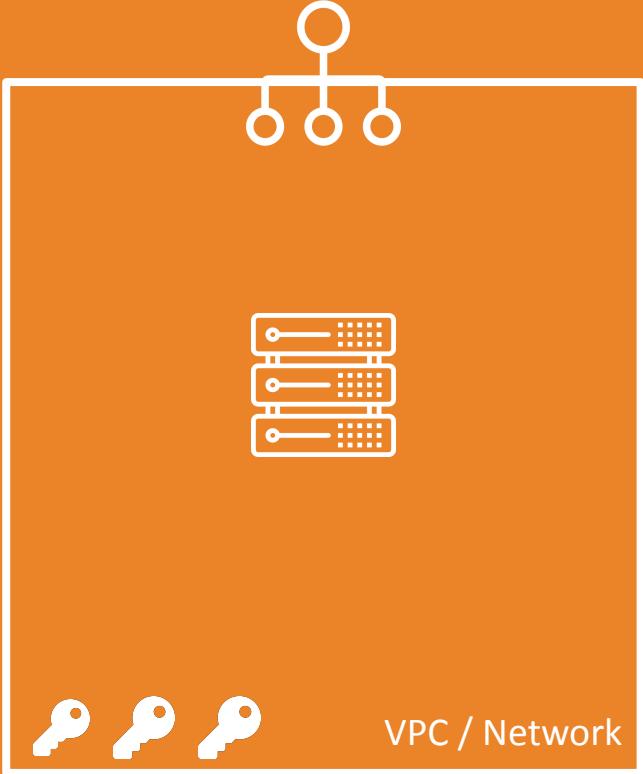


```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
kind: OpenStackMachine
```

```
spec:
  identityRef:
    kind: Secret
    name: cloud-config
```

```
flavor: n1.medium
image: ubuntu-2004-kube-v1.22.8
rootVolume:
  diskSize: 60
```

***OK. We have a machine
now. Let's talk about
how to bootstrap it.***



Kind: KubeADMConfig

- **Managed by “capi-kubeadm-bootstrap-controller-manager” (Bootstrap Operator)**
- **CR to declare**
 - **Bootstrap configuration like kubeadm config**
 - **Files to mount**
 - **Users to create**
 - **Commands to run**

kind: KubeadmConfig : Spec



As a Cluster API user

I want to bootstrap my machine

By injecting these files

By using this join configuration

By using this kubelet configuration

By running these commands

By using this format

```
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmConfig

spec:
  files:
    - content: |
        ssh-ed25519 AAAAC3NzaC1IzDI1NTE5AA...
        path: /etc/ssh/trusted-user-ca-keys.pem
        permissions: "0600"
    joinConfiguration:
      discovery:
        bootstrapToken:
          apiServerEndpoint: xxx.yyy.ttt.zzz:6443
          token: u7zw64.mytoken
        nodeRegistration:
          kubeletExtraArgs:
            eviction-hard: memory.available<200Mi
            feature-gates: ExpandPersistentVolumes=true
            name: '{{ local_hostname }}'
      postKubeadmCommands:
        - systemctl restart sshd
  format: cloud-config
```

How the bootstrapping works - I

Bootstrap operator creates a secret based on the spec in the KubeAdmConfig CR.

```
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmConfig
status:
dataSecretName: demo1-region1-8342e0a3-lwlqx
ready: true
```

How the bootstrapping works - 2

The secret content is in one of the supported formats.

cloud-config files are special scripts designed to be run by the cloud-init service.

```
## template: jinja
#cloud-config

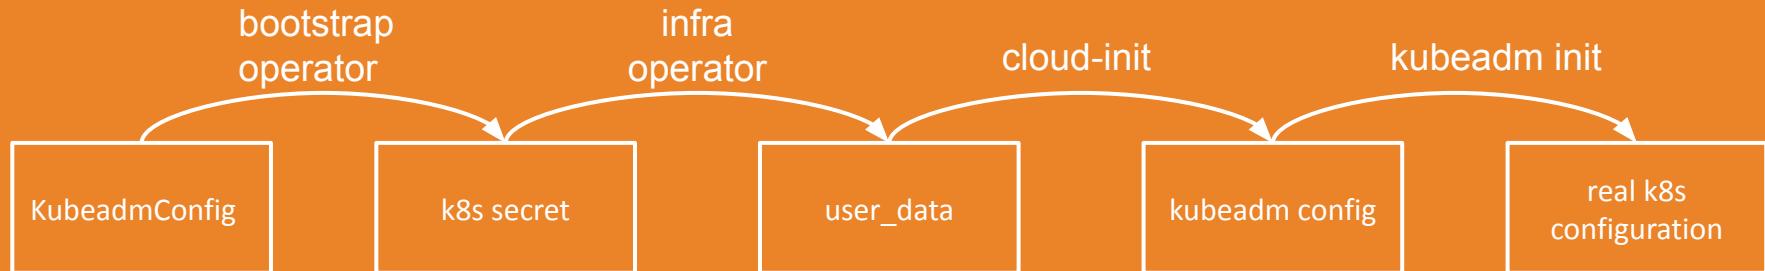
write_files:
- path: /etc/ssh/trusted-user-ca-keys.pem
  permissions: '0600'
  content: |
    ssh-ed25519 dfdsfs
    vault-ca@vault.operations.giantswarm.io

- path: /run/kubeadm/kubeadm-join-config.yaml
  owner: root:root
  permissions: '0640'
  content: |
  ---
  apiVersion: kubeadm.k8s.io/v1beta3

runcmd:
- kubeadm join --config
  /run/kubeadm/kubeadm-join-config.yaml
  - "systemctl restart sshd"
users:
- name: giantswarm
  sudo: ALL=(ALL) NOPASSWD:ALL
```

How the bootstrapping works - 3

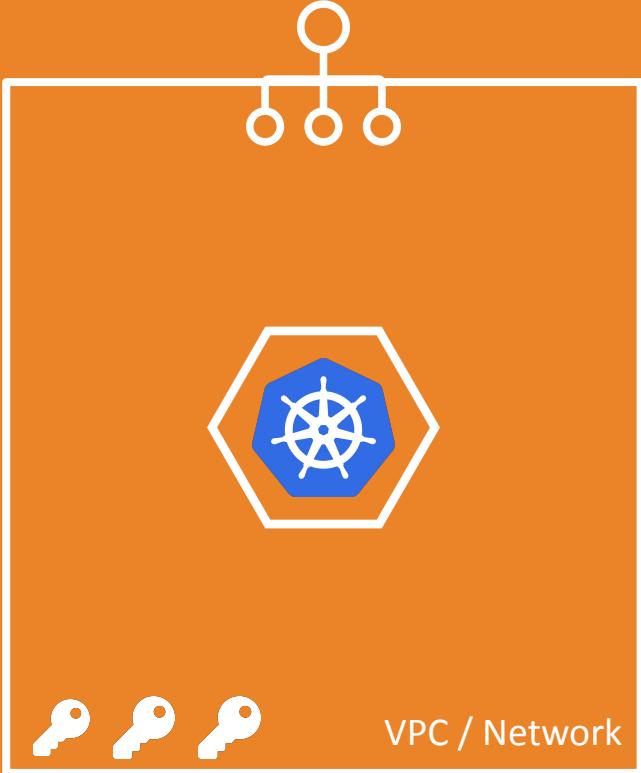
Infra operator fetches bootstrap configuration from Machine CR (by following owner references) and uses that as user_data for cloud-init while creating a machine.



**We converted the
machine to a k8s node,
which is great!**

We plan to create more.

**Is there a way to
manage them as a set
instead of one by one?**



Cluster API follows k8s approach!



Deployment



Machine Deployment



Replica Set



Machine Set



Pod



Machine

kind: MachineDeployment

- **Managed by “capi-controller-manager” (Core Operator)**
- **Higher level CR to manage Machines like Deployment in Kubernetes**
- **Refers update strategies and update configuration**

kind: MachineDeployment : Spec

part of core api

like k8s deployment

update strategy

machine spec

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: MachineDeployment

spec:
  minReadySeconds: 0
  progressDeadlineSeconds: 600
  replicas: 5
  selector:
    matchLabels:
      cluster.x-k8s.io/cluster-name: demo1
      cluster.x-k8s.io/deployment-name: demo1-region1

  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
      type: RollingUpdate

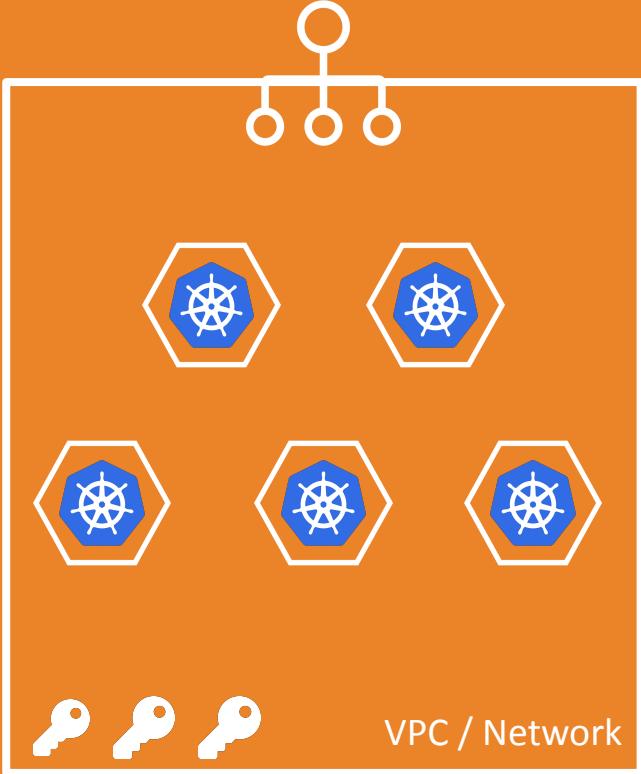
  template:
    MACHINE_TEMPLATE
```

Wait! To be able to create all these nodes, we need a control plane to join.

Where is the control plane?

In the beginning, in the Cluster CR, we mentioned KubeadmControlPlane.

Let's talk about it.



kind: KubeADMControlPlane

- **Managed by “capi-kubeadm-control-plane-controller-manager” (ControlPlane Operator)**
- **Referred by Cluster**
- **Refers control-plane configuration**
 - **api-server**
 - **controller-manager**
 - **scheduler**
 - **etcd**
- **Contains fields of KubeADMConfig too since every control plane node is also a node :)**

kind: KubeADMControlPlane : Spec - I

Control-plane configuration

```
apiVersion: controlplane.cluster.x-k8s.io/v1beta1
kind: KubeADMControlPlane

spec:
  kubeADMConfigSpec:
    clusterConfiguration:
      apiServer:
        extraArgs:
          enable-admission-plugins: NamespaceLifecycle,LimitRanger...
      controllerManager:
        extraArgs:
          authorization-always-allow-paths: /healthz,/readyz,/livez,/metrics
          bind-address: 0.0.0.0
      etcd:
        local:
          extraArgs:
            listen-metrics-urls: http://0.0.0.0:2381
          imageRepository: giantswarm
          imageTag: 3.5.4-0-k8s
      scheduler:
```

kind: KubeADMControlPlane : Spec - 2

- Like KubeADMConfig since all control-plane nodes are also nodes!**
- Refers InfraMachineTemplate CR**
- Like MachineDeployment but for control-plane machines**

```
apiVersion: controlplane.cluster.x-k8s.io/v1beta1
kind: KubeADMControlPlane
spec:
  kubeADMConfigSpec:
    files:
      format: cloud-config
    initConfiguration:
    joinConfiguration:
    preKubeADMCommands:
    machineTemplate:
      infrastructureRef:
        apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
        kind: OpenStackMachineTemplate
        name: demo1-control-plane-91da4cd9
        namespace: org-multi-project
      replicas: 3
      rolloutStrategy:
        rollingUpdate:
          maxSurge: 1
        type: RollingUpdate
```

How control plane provisioning works - I

```
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmControlPlane
metadata:
  name: demo1-rrsnf
spec:
  kubeadmConfigSpec:
    clusterConfiguration:
      apiServer:
        certSANs:
          - 127.0.0.1
          - localhost
          - api.demo1.test.erkan.io
```

}



```
kind: Secret
metadata:
  name: demo1-ca
```

```
kind: Secret
metadata:
  name: demo1-etcd
```

```
kind: Secret
metadata:
  name: demo1-sa
```

```
kind: Secret
metadata:
  name: demo1-proxy
```

```
kind: Secret
metadata:
  name: demo1-kubeconfig
```

When cluster infrastructure is ready, ControlPlane operator provisions certificates and kubeconfig for the cluster. It puts those into k8s secrets.

Everyone uses this to access WC.

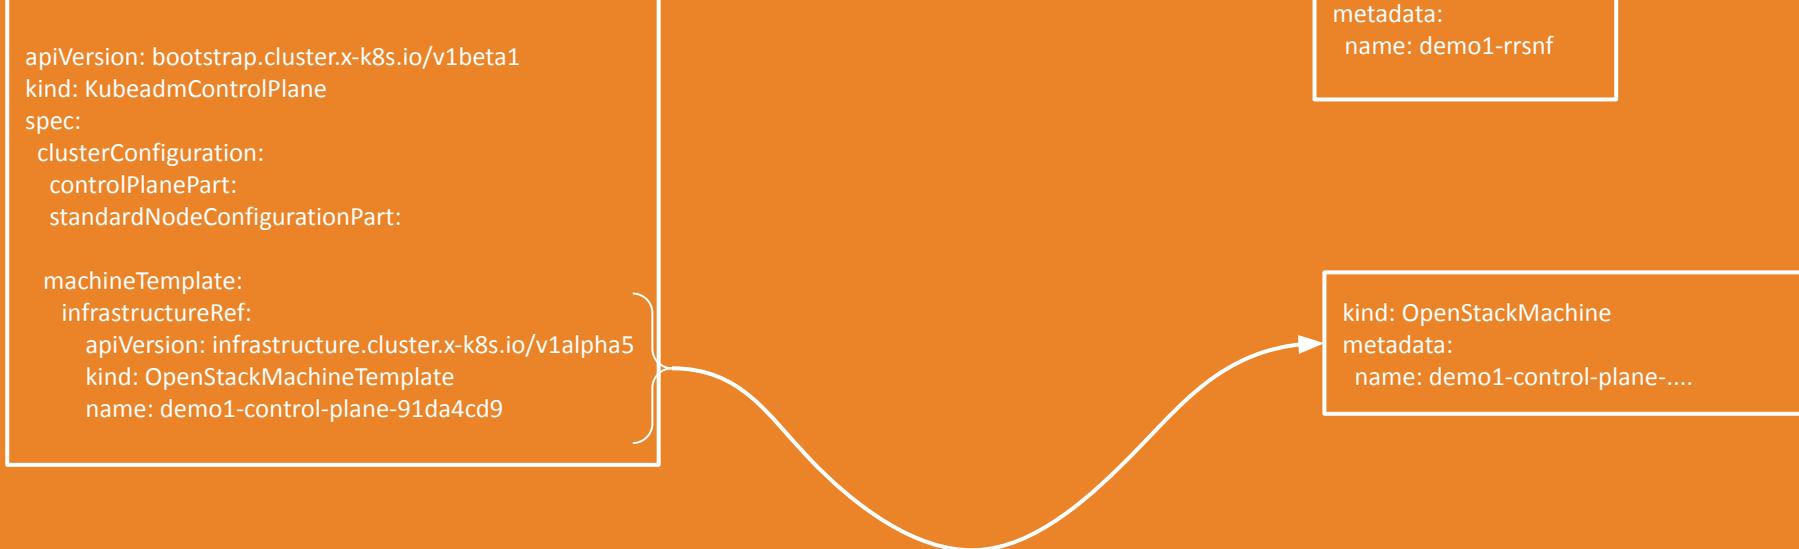
How control plane provisioning works - 2

```
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeADMControlPlane
metadata:
  name: demo1-rrsnf
spec:
  clusterConfiguration:
    controlPlanePart:
    standardNodeConfigurationPart:
      machineTemplate:
        infrastructureRef:
          apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
          kind: OpenStackMachineTemplate
          name: demo1-control-plane-91da4cd9
```



ControlPlane operator creates a special KubeADMConfig for controlplane nodes by using control-plane specific and generic configuration in KubeADMControlPlane CR. This KubeADMConfig includes certificates too.

How control plane provisioning works - 3

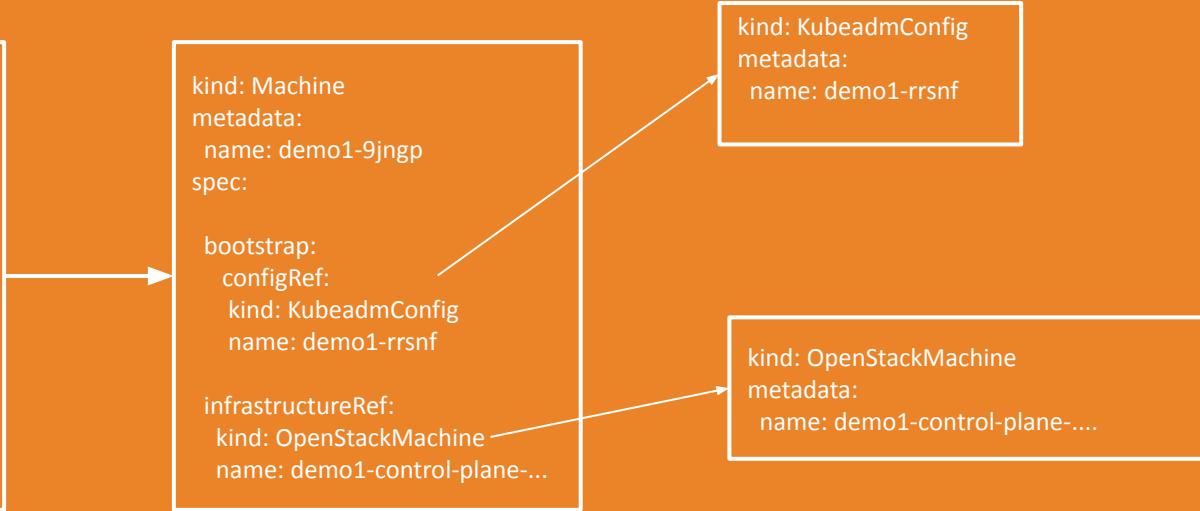


ControlPlane operator creates a InfraMachine CR by using InfraMachineTemplate reference in KubeadmControlPlane CR.

How control plane provisioning works - 4

```
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmControlPlane
spec:
  clusterConfiguration:
    controlPlanePart:
      standardNodeConfigurationPart:

  machineTemplate:
    infrastructureRef:
      apiVersion: infrastructure.cluster.x-k8s.io/v1alpha5
      kind: OpenStackMachineTemplate
      name: demo1-control-plane-91da4cd9
```

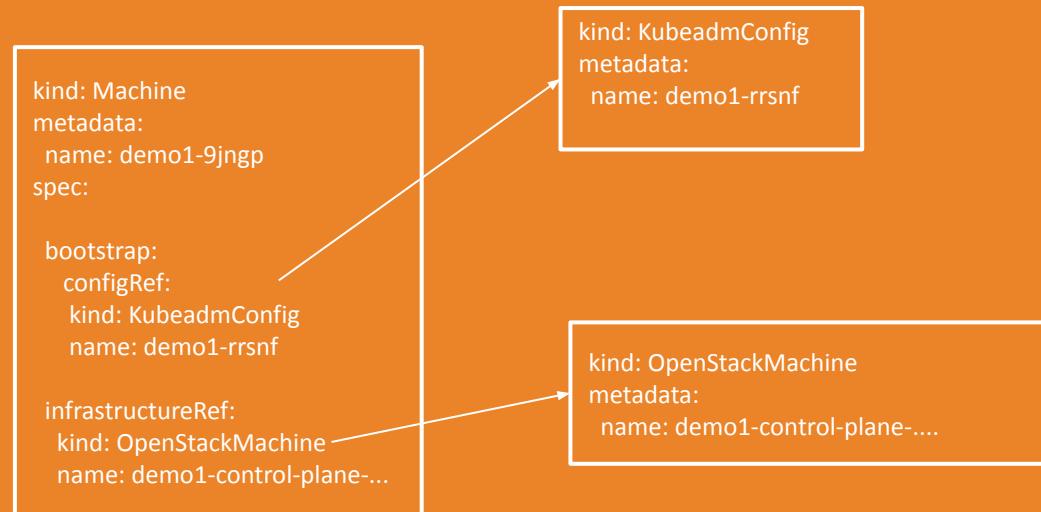


ControlPlane operator finally creates a Machine CR that refers the KubeadmConfig and InfraMachine CRs created earlier.

How control plane provisioning works - 5

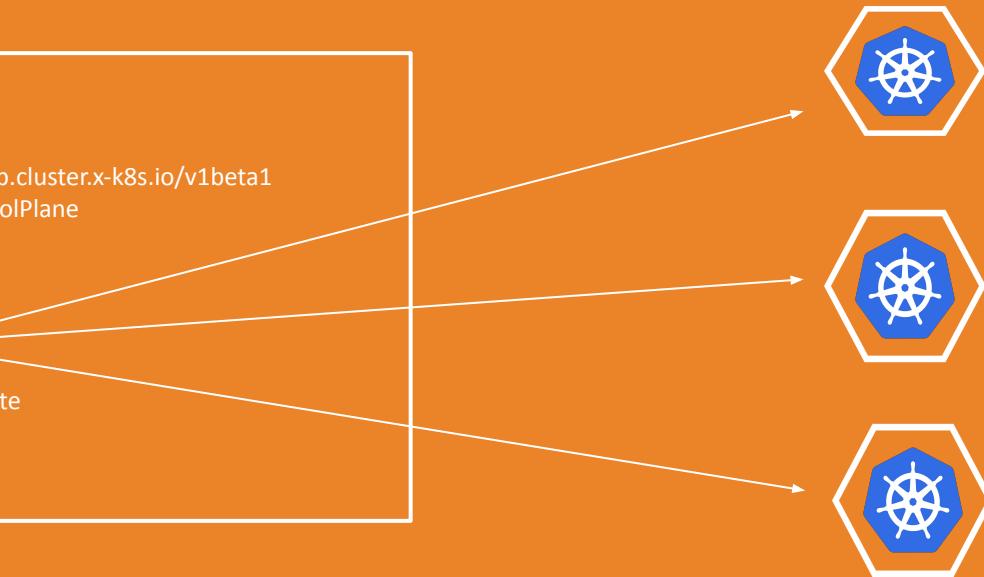
**As we talked earlier,
Infra operator creates
a machine by using
cloud-config provided
by Bootstrap operator.**

**At the end, the created
node is a
control-plane node.**



How control plane provisioning works - 6

```
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmControlPlane
spec:
  replicas: 3
  rolloutStrategy:
    rollingUpdate:
      maxSurge: 1
      type: RollingUpdate
```



ControlPlane operator provisions new control-plane nodes according to “replicas” field and rollout nodes during upgrades according to KubeadmControlPlane spec.

How control plane provisioning works - 7

- **ControlPlane operator is not only creating machines like MachineSet.**
- **It also access & manipulate workload cluster internals to manage control plane nodes in a stable way.**

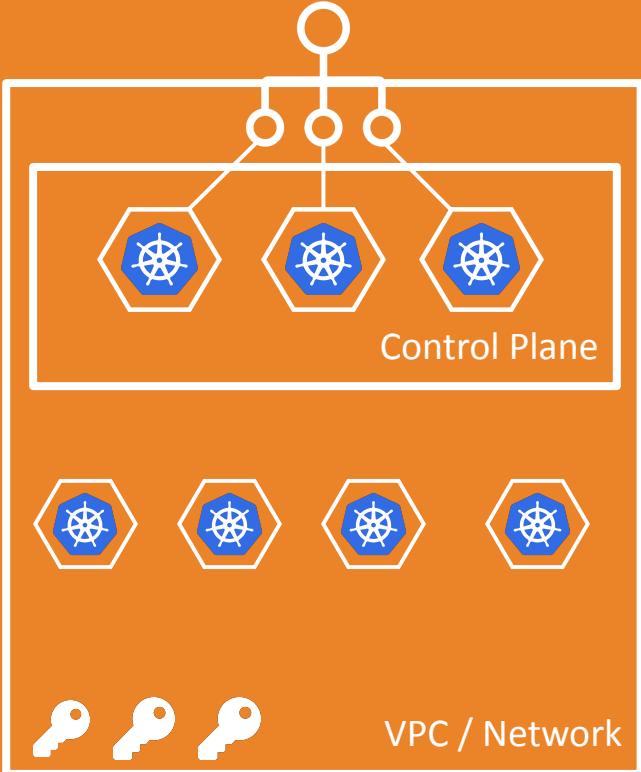
How control plane provisioning works - 8

```
type WorkloadCluster interface {
    // Basic health and status checks.
    ClusterStatus(ctx context.Context) (ClusterStatus, error)
    UpdateStaticPodConditions(ctx context.Context, controlPlane *ControlPlane)
    UpdateEtcdConditions(ctx context.Context, controlPlane *ControlPlane)
    EtcMembers(ctx context.Context) ([]string, error)
    GetAPIServerCertificateExpiry(ctx context.Context, kubeadmConfig *bootstrapv1.KubeadmConfig, nodeName string) (*time.Time, error)

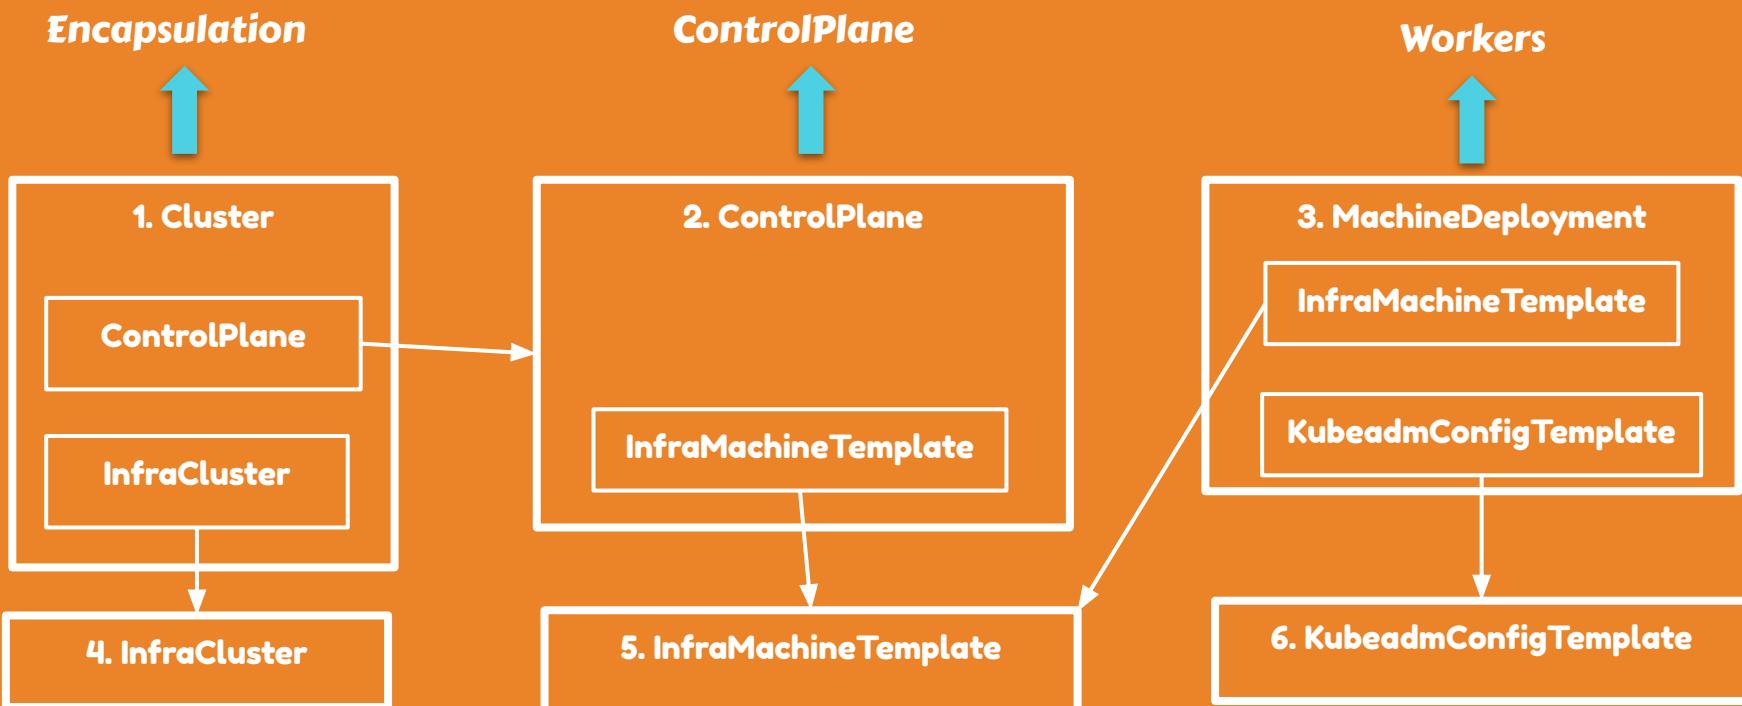
    // Upgrade related tasks.
    ReconcileKubeletRBACBinding(ctx context.Context, version semver.Version) error
    ReconcileKubeletRBACRole(ctx context.Context, version semver.Version) error
    UpdateKubernetesVersionInKubeadmConfigMap(ctx context.Context, version semver.Version) error
    UpdateImageRepositoryInKubeadmConfigMap(ctx context.Context, imageRepository string, version semver.Version) error
    UpdateEtcdVersionInKubeadmConfigMap(ctx context.Context, imageRepository, imageTag string, version semver.Version) error
    UpdateEtcdExtraArgsInKubeadmConfigMap(ctx context.Context, extraArgs map[string]string, version semver.Version) error
    UpdateAPIServerInKubeadmConfigMap(ctx context.Context, apiServer bootstrapv1.APIServer, version semver.Version) error
    UpdateControllerManagerInKubeadmConfigMap(ctx context.Context, controllerManager bootstrapv1.ControlPlaneComponent, version semver.Version) error
    UpdateSchedulerInKubeadmConfigMap(ctx context.Context, scheduler bootstrapv1.ControlPlaneComponent, version semver.Version) error
    UpdateKubeletConfigMap(ctx context.Context, version semver.Version) error
    UpdateKubeProxyImageInfo(ctx context.Context, kcp *controlplanev1.KubeadmControlPlane, version semver.Version) error
    UpdateCoreDNS(ctx context.Context, kcp *controlplanev1.KubeadmControlPlane, version semver.Version) error
    RemoveEtcdMemberForMachine(ctx context.Context, machine *clusterv1.Machine) error
    RemoveMachineFromKubeadmConfigMap(ctx context.Context, machine *clusterv1.Machine, version semver.Version) error
    RemoveNodeFromKubeadmConfigMap(ctx context.Context, nodeName string, version semver.Version) error
    ForwardEtcdLeadership(ctx context.Context, machine *clusterv1.Machine, leaderCandidate *clusterv1.Machine) error
    AllowBootstrapTokensToGetNodes(ctx context.Context) error

    // State recovery tasks.
    ReconcileEtcdMembers(ctx context.Context, nodeNames []string, version semver.Version) ([]string, error)
}
```

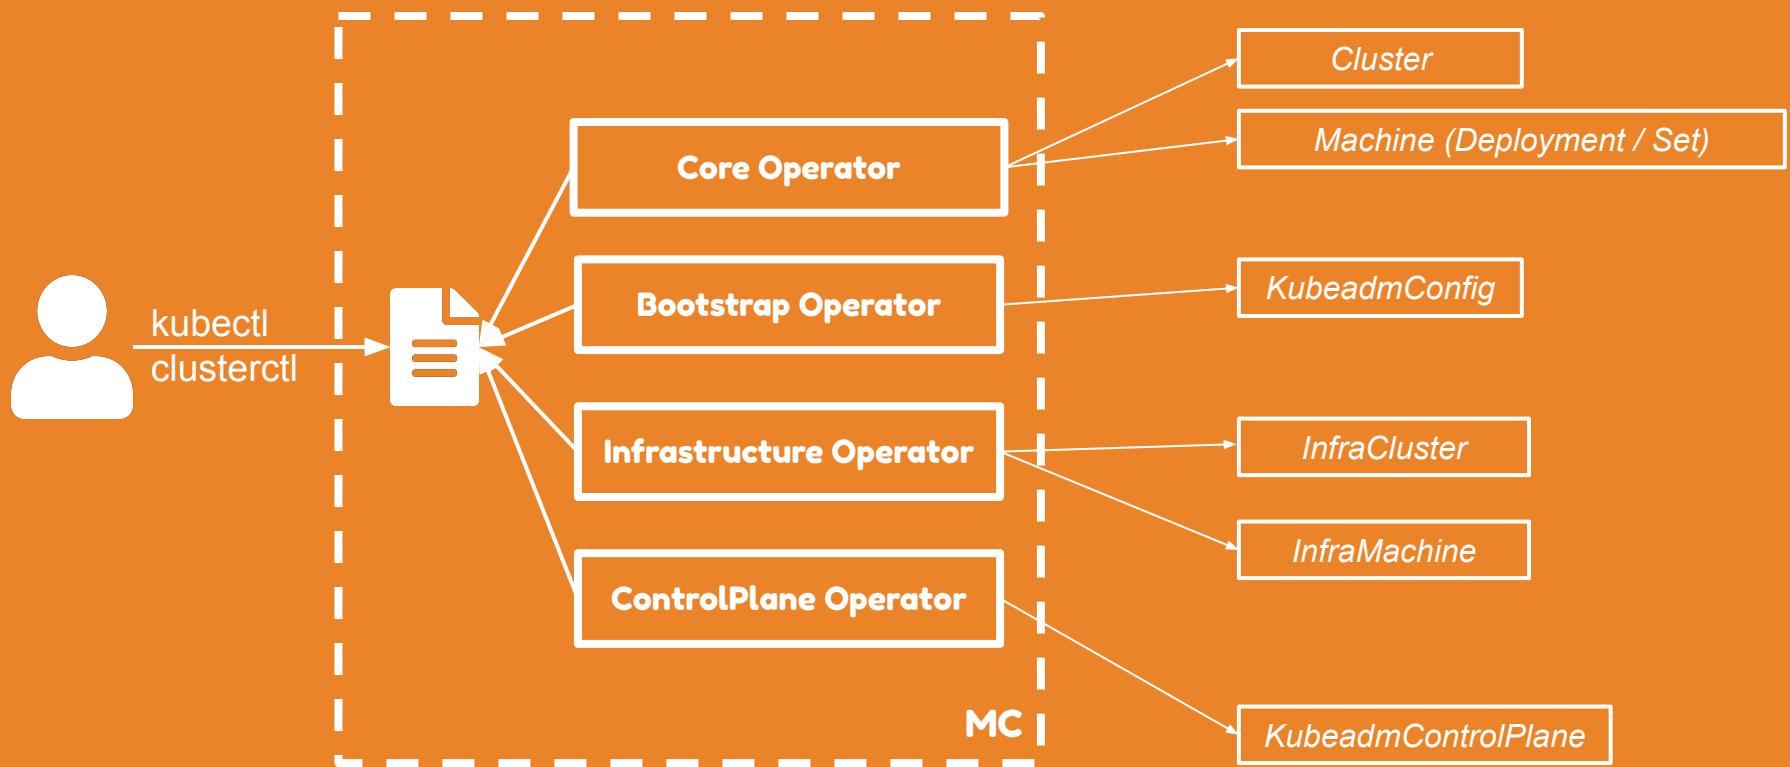
**We have a *fully working*
cluster now.**



RECAP-I



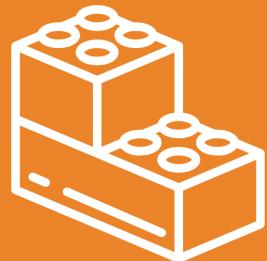
RECAP - 2



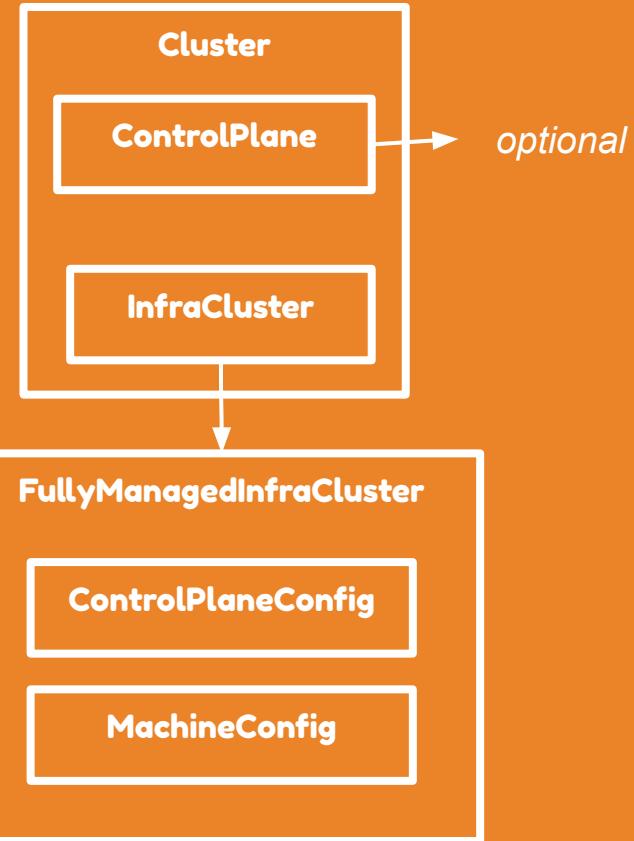
RECAP - 3

```
➔ clusterctl describe cluster demo1 --grouping=false --show-machinesets  
NAME  
Cluster/demo1  
  ClusterInfrastructure - OpenStackCluster/demo1  
  ControlPlane - KubeadmControlPlane/demo1  
    Machine/demo1-9jngp  
      MachineInfrastructure - OpenStackMachine/demo1-control-plane-91da4cd9-ggbhd  
  Workers  
    MachineDeployment/demo1-region1  
    MachineSet/demo1-region1-7776fdbf86  
      Machine/demo1-region1-7776fdbf86-d981v  
        MachineInfrastructure - OpenStackMachine/demo1-default-91da4cd9-6q9km  
      Machine/demo1-region1-7776fdbf86-j5r8n  
        MachineInfrastructure - OpenStackMachine/demo1-default-91da4cd9-7hvwc  
      Machine/demo1-region1-7776fdbf86-t9jz2  
        MachineInfrastructure - OpenStackMachine/demo1-default-91da4cd9-m7gn2
```

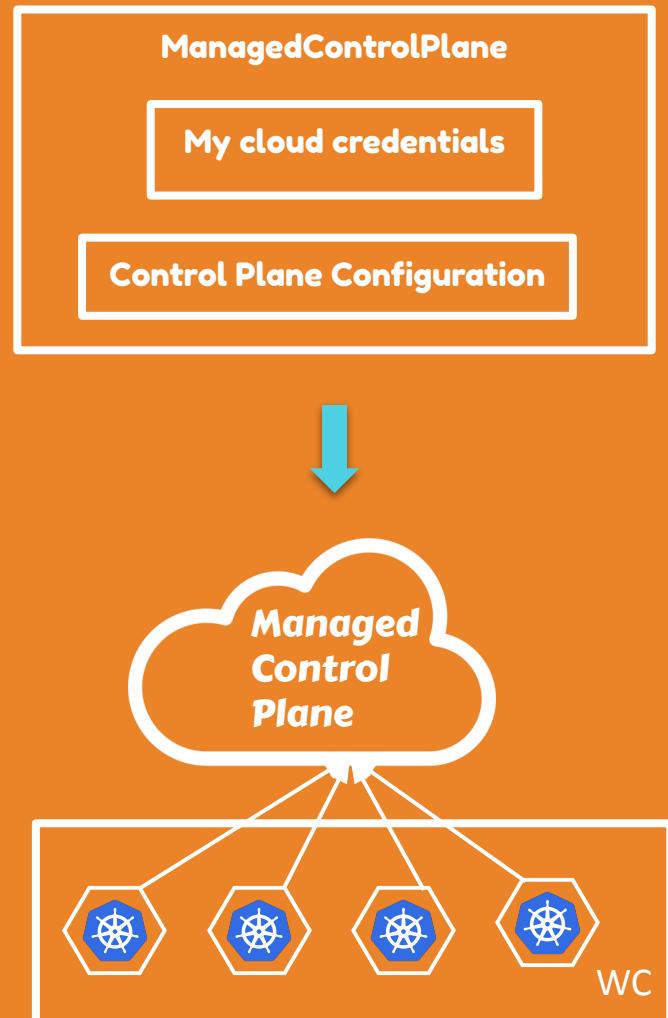
Let's talk about interchangeability again.



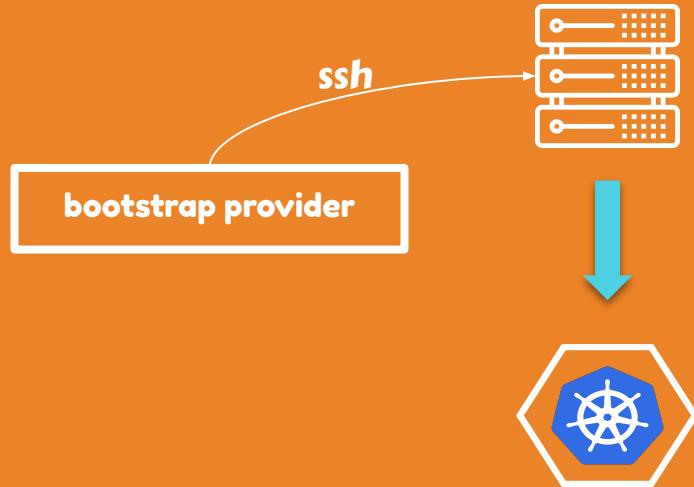
An infra provider can implement only InfraCluster without using others CRs. It is up to infra provider.



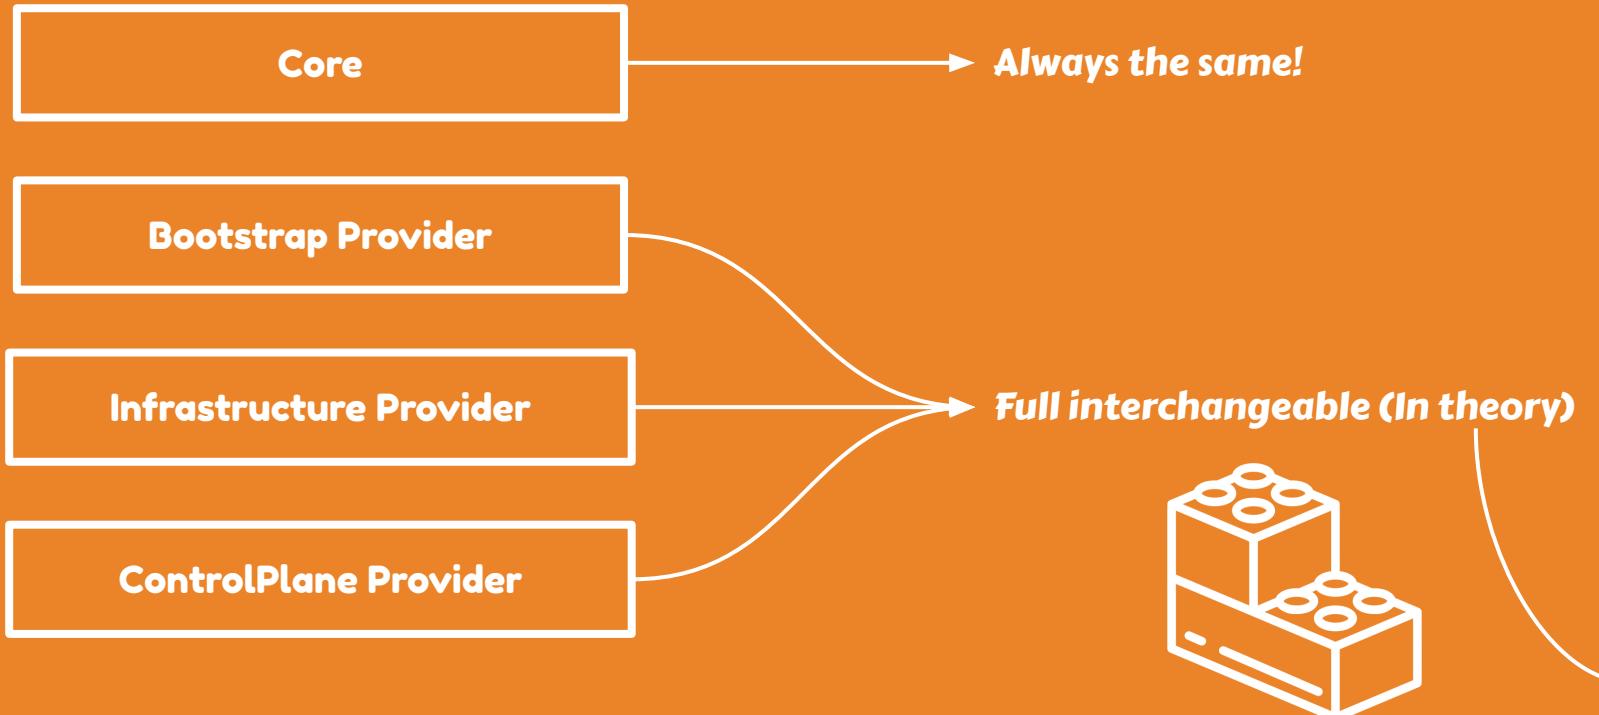
A controlplane provider can provision the control plane in a cloud service without using bootstrap/machine apis.



A bootstrap provider can bootstrap a machine by connecting it to through ssh and running commands without using cloud-init. It is up to bootstrap provider.



Cluster API components are interchangeable!



BONUS PART



Generated by DALL-E

kind: MachineHealthCheck

- **Part of core api.**
- **Core operator watches machine. If there is a machine which is not in Ready state, it deletes that machine and create a new one.**

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: MachineHealthCheck
metadata:
  name: demo1
spec:
  clusterName: demo1
  maxUnhealthy: 40%
  nodeStartupTimeout: 20m0s

  selector:
    matchLabels:
      cluster.x-k8s.io/cluster-name: demo1

  unhealthyConditions:
    - status: Unknown
      timeout: 10m0s
      type: Ready
    - status: "False"
      timeout: 10m0s
      type: Ready
```

kind: ClusterResourceSet



As a Cluster API user

In every cluster with these labels

**I want to create this k8s object
automatically.**

```
apiVersion: addons.cluster.x-k8s.io/v1beta1
kind: ClusterResourceSet

metadata:
  name: demo1-coredns

spec:
  clusterSelector:
    matchLabels:
      cluster.x-k8s.io/cluster-name: demo1

resources:
  - kind: ConfigMap
    name: demo1-coredns
    strategy: ApplyOnce
```

kind: ClusterClass - I

- **Experimental feature**
- **Feature gate name:**
ClusterTopology
- **Creating a topology once
and using many times**

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: ClusterClass
metadata:
  name: openstack-clusterclass-v0.1.0
spec:
  controlPlane:
    ref:
      apiVersion: controlplane.cluster.x-k8s.io/v1beta1
      kind: KubeadmControlPlaneTemplate
      name: openstack-clusterclass-v0.1.0
      namespace: default
  machineInfrastructure:
    ref:
      kind: openstackMachineTemplate
      apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
      name: openstack-clusterclass-v0.1.0
      namespace: default
  infrastructure:
    ref:
      apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
      kind: OpenstackClusterTemplate
      name: openstack-clusterclass-v0.1.0-control-plane
      namespace: default
  workers:
    machineDeployments:
      - class: default-worker
        template:
          bootstrap:
            ref:
              apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
              kind: KubeadmConfigTemplate
              name: Openstack-clusterclass-v0.1.0-default-worker
              namespace: default
        infrastructure:
          ref:
            apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
            kind: OpenstackMachineTemplate
            name: openstack-clusterclass-v0.1.0-default-worker
            namespace: default
```

kind: ClusterClass - 2

- **Experimental feature**
- **Feature gate name:**
ClusterTopology
- **Creating a topology once
and using many times**

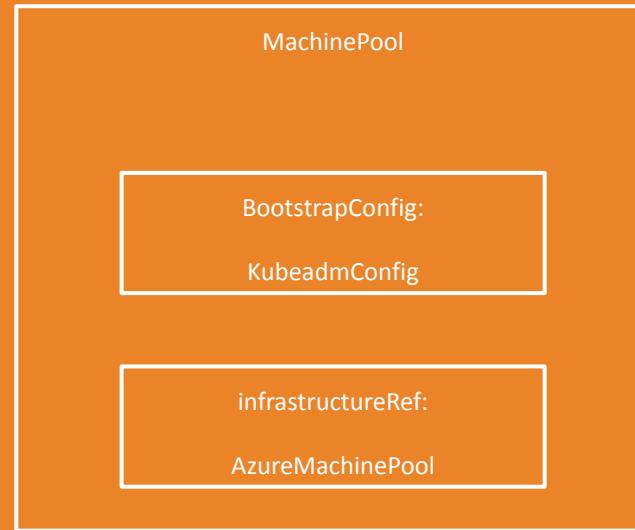
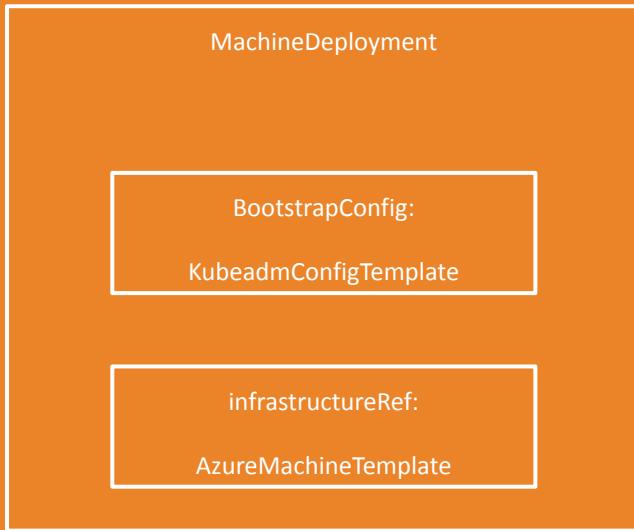


```
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
metadata:
  name: my-openstack-cluster
spec:
  topology:
    class: openstack-clusterclass
    version: v1.22.4
    controlPlane:
      replicas: 3
      metadata:
        labels:
          cpLabel: cpLabelValue
        annotations:
          cpAnnotation: cpAnnotationValue
    workers:
      machineDeployments:
        - class: default-worker
          name: md-0
          replicas: 4
          metadata:
            labels:
              mdLabel: mdLabelValue
            annotations:
              mdAnnotation: mdAnnotationValue
          failureDomain: region
```

kind: MachinePool - I

- **Experimental feature**
- **Feature gate name: MachinePool**
- **Like MachineDeployment but doesn't use MachineSet**
- **Delegates the responsibility of these concerns to an infrastructure provider specific resource such as AWS Auto Scale Groups, GCP Managed Instance Groups, and Azure Virtual Machine Scale Sets.**

kind: MachinePool - 2



**Let me talk about my company
for 5 minutes**

**But not like
a salesman.
It will be
like tips and
tricks.**



Sauce of

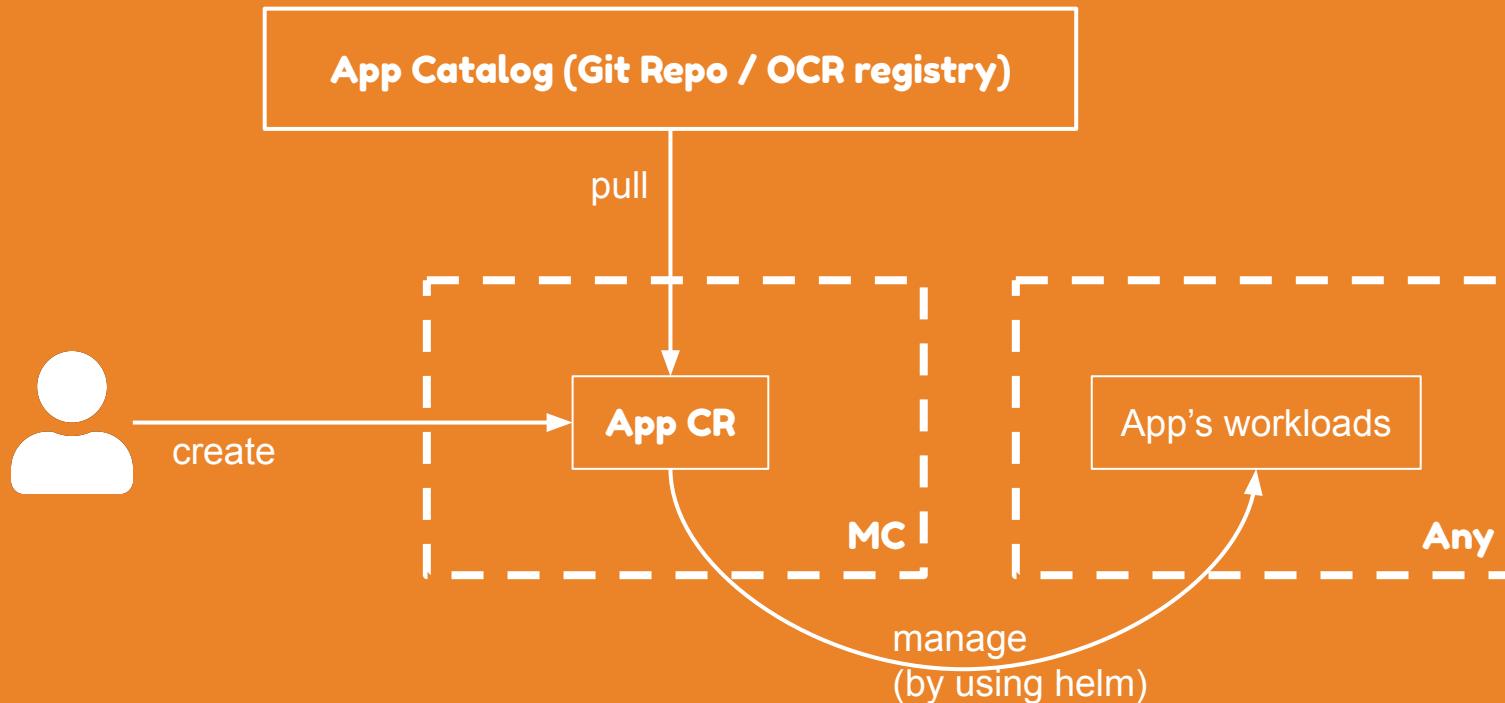


Giant Swarm



PUBLIC

App Platform - I



App Platform - 2

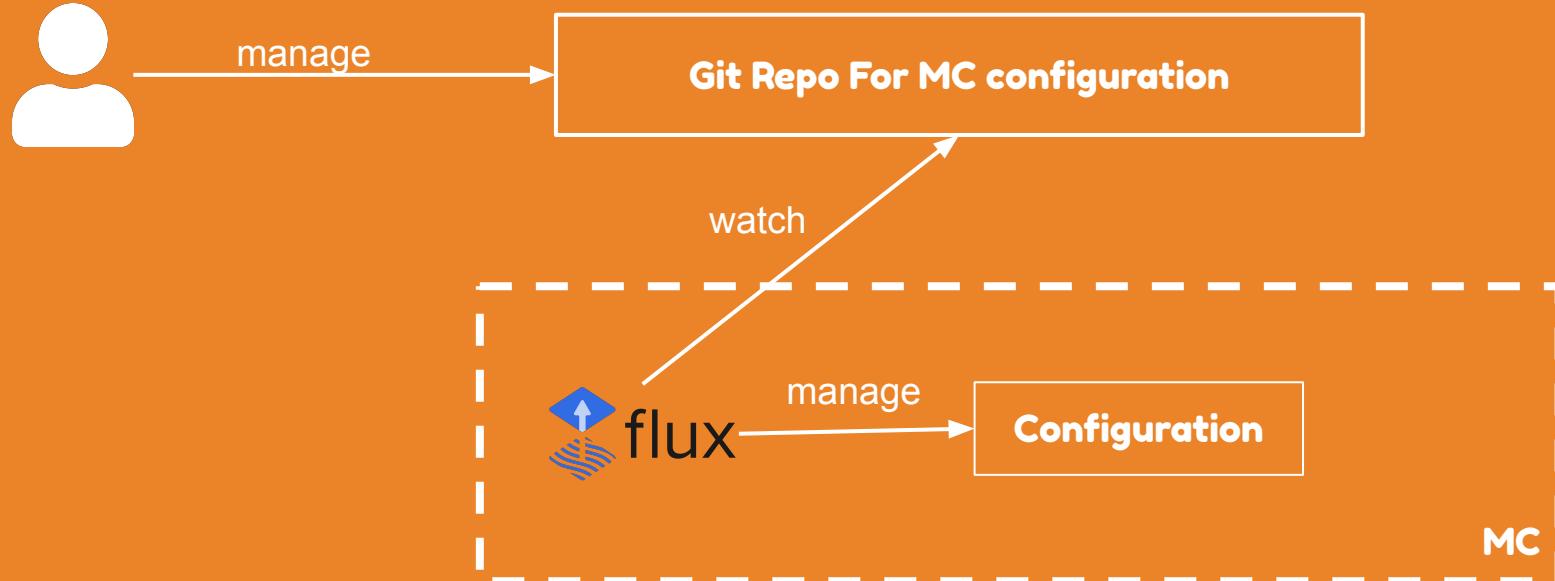
- **There are lots of apps in the app catalogs**
 - [giantswarm/giantswarm-catalog](#)
 - [giantswarm/default-catalog](#)
- **We use the app platform to create WCs too.**
 - [giantswarm/cluster-api-app](#)
 - [giantswarm/cluster-catalog](#)
- **Upgrade = changing app version**

```
apiVersion: application.giantswarm.io/v1alpha1
kind: App
metadata:
  name: demo1-cluster
  namespace: org-multi-project
spec:
  catalog: cluster
  name: cluster-openstack
  namespace: org-multi-project
  userConfig:
    configMap:
      name: demo1-cluster-userconfig
      namespace: org-multi-project
  version: 0.16.1
```

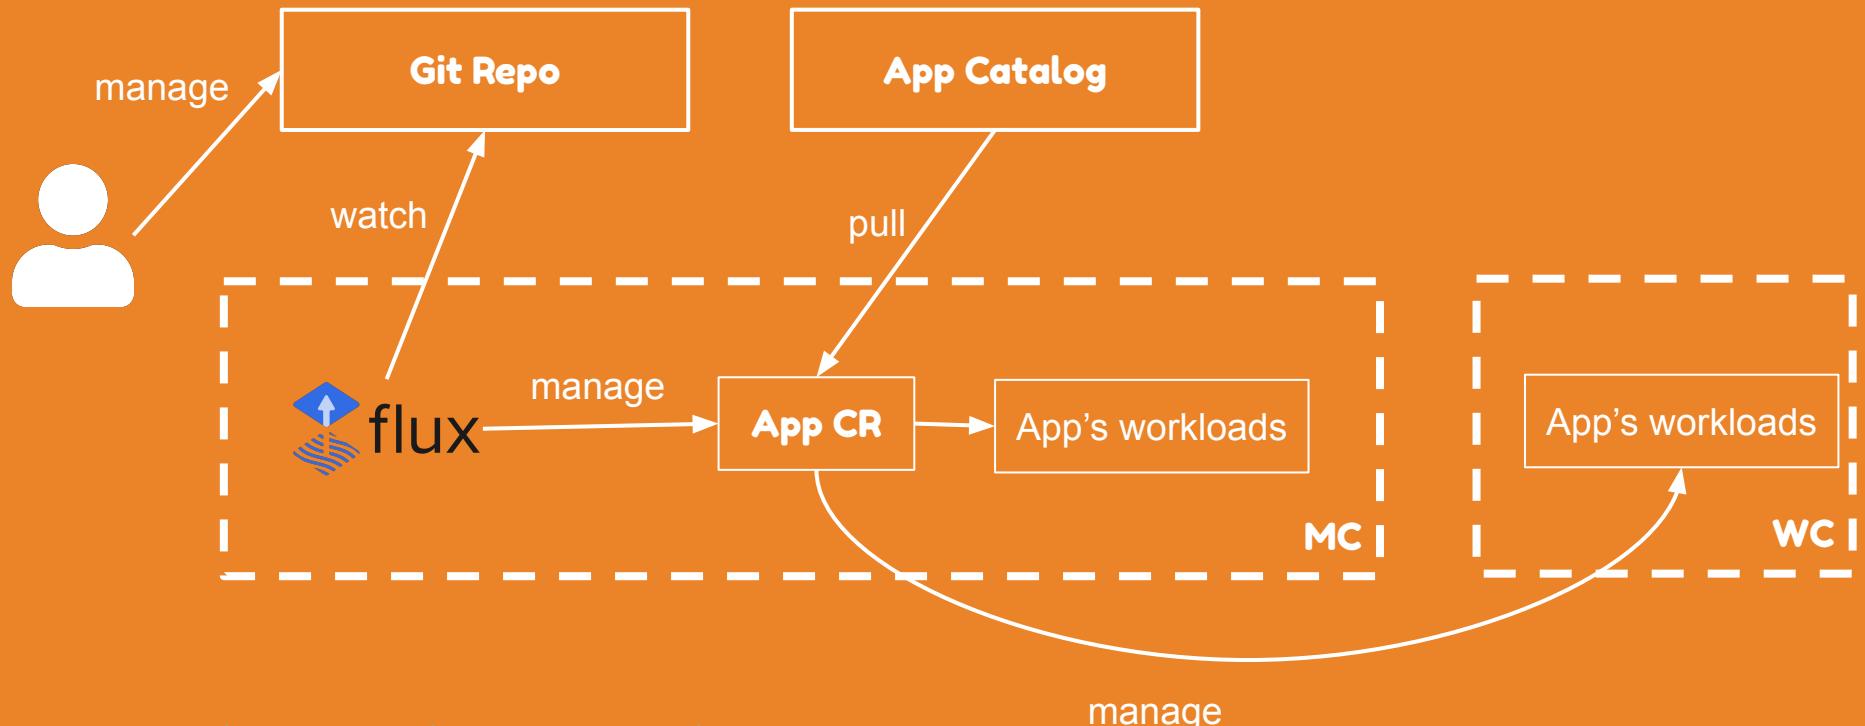
Example Apps

- **Observability**
 - **Prometheus** : You can access any WC by just changing a URL param
 - **Grafana**: There are many built-in dashboard
 - **Loki, EFK stack**
 - **Note: GiantSwarm monitors all clusters and gives 7/24 support.**
- **Security**
 - **Kyverno**: There are some built-in policies for secure clusters
 - **Falco**
 - **Trivy**
- **Connectivity**
 - **linkerd**
 - **Kong, ingress**

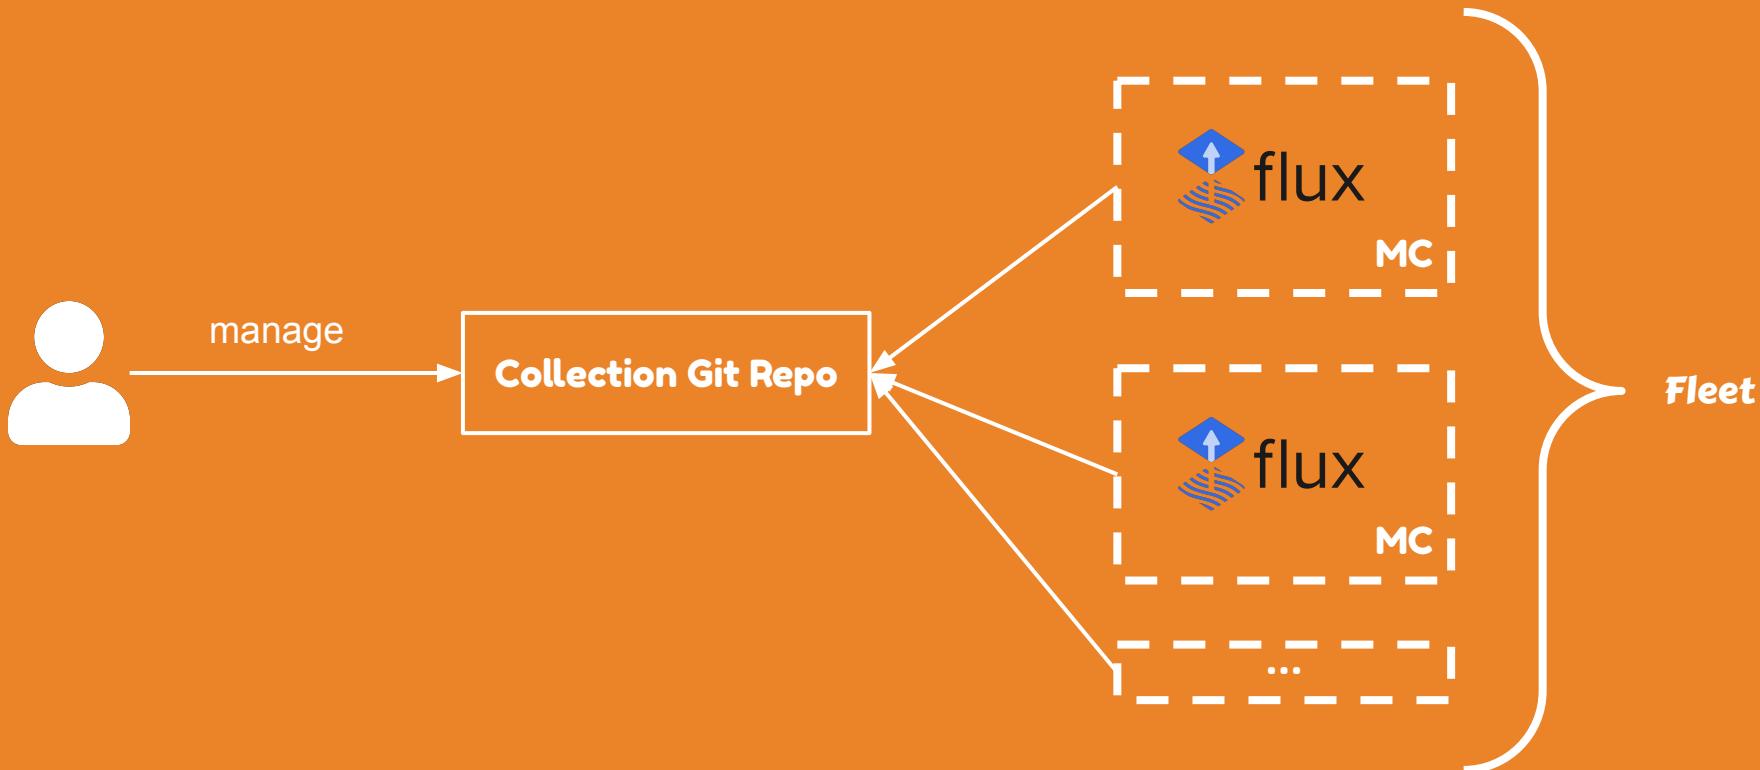
Gitops + Flux - I



Gitops + Flux - 2



Gitops + Flux - 3: Fleets & Collections



e.g. <https://github.com/giantswarm/openstack-app-collection>

Management Tools

- **Powerful internal admin tools**
 - **Lists management clusters**
 - **Get kubeconfig for MC/WC**
 - **Get access to MC/WC**
 - **Daily ops activities (deploy, rollout, drain etc.)**
 - **Open MC apps like prometheus, grafana etc.**
- **We have powerful user tools like kubectl-gs**
 - **Easy templating**
 - **WC management, especially GS types (e.g. apps)**
- **Web UI (giantswarm/happa)**



User Experience

Open source tools in upstream repositories are good but not enough for a smooth user experience. We

- **implement additional/helper operators**
 - **for easy upgrades with declarative approach like GitOps**
 - **for backup**
 - **for cleanup**
- **create useful helm charts**
(See `cluster-\$provider` repos in giantswarm.org)





The problem domain is huge and complex.

Ops is inevitable.

Support is must.

THANK
YOU!



for listening!